




© Raquel H. Ribeiro

Exposing fraud with machine learning

DS Portugal Meetup @ UPtec
Porto, 24th May 2017



Roadmap

1. A payments world overview
 2. Where there is value, there is fraud
 3. Using machine learning to uncover and prevent fraud
- 



About Feedzai

- Young start-up with a Science & Engineering DNA
- Expertise in forecasting and fraud detection
- Processes transactions worth more than Portugal's GDP

About DS team

- Currently over 25 data scientists from various backgrounds in Lisbon, Coimbra, Porto, New York and Atlanta - yet still growing!
- Delivery team works in PoCs and established projects, from banking to e-commerce
- Research team leads innovation
- Code review and reading groups are standard practice

The payments world

the customer

the merchant

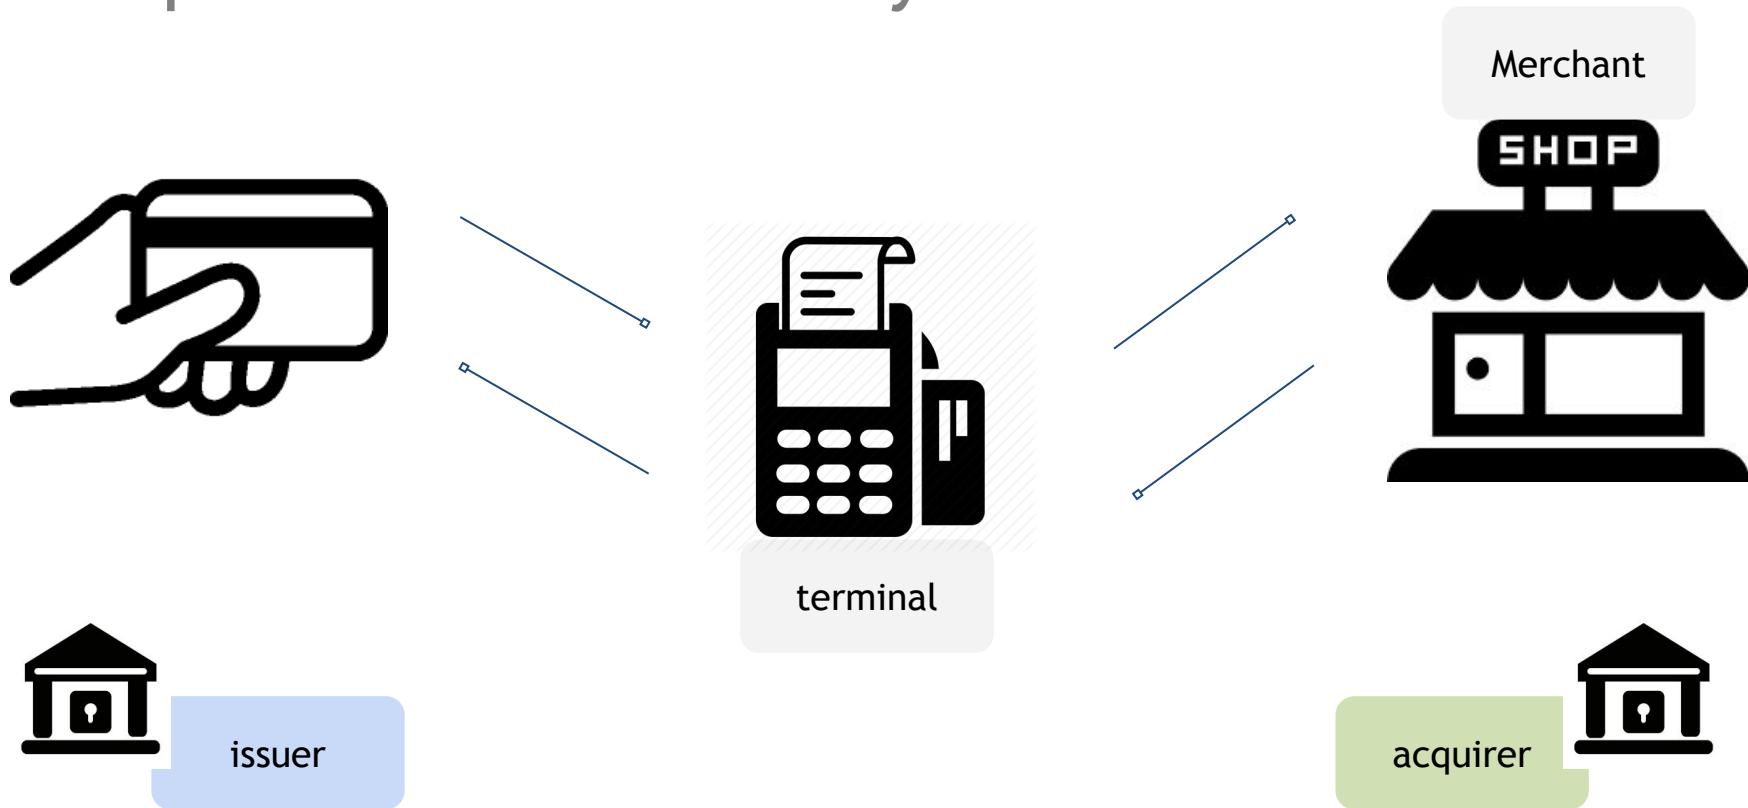
the acquirer

the issuer

How to buy services and goods?



In a point of sale near you




When you swipe your card



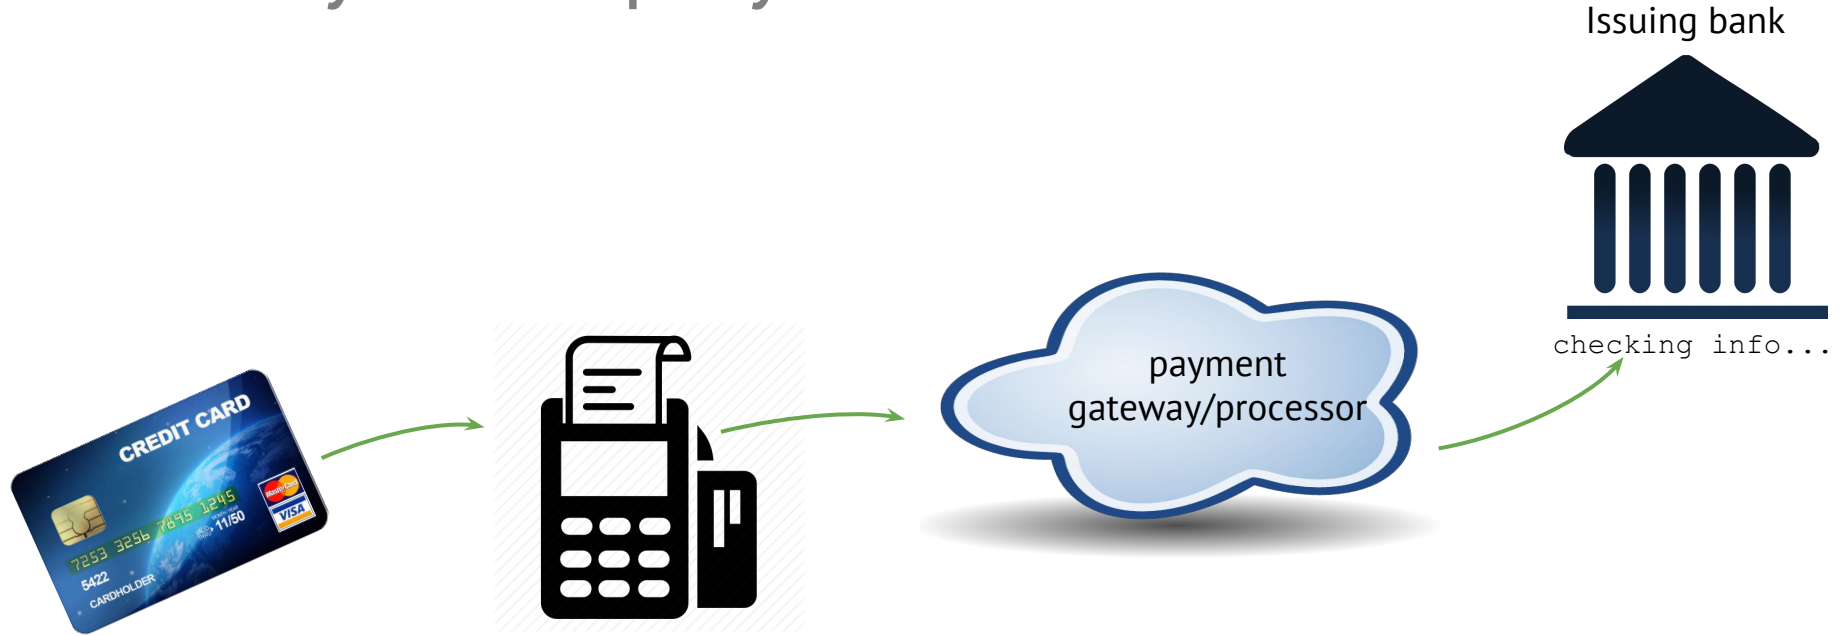
When you swipe your card



When you swipe your card



When you swipe your card



When you swipe your card



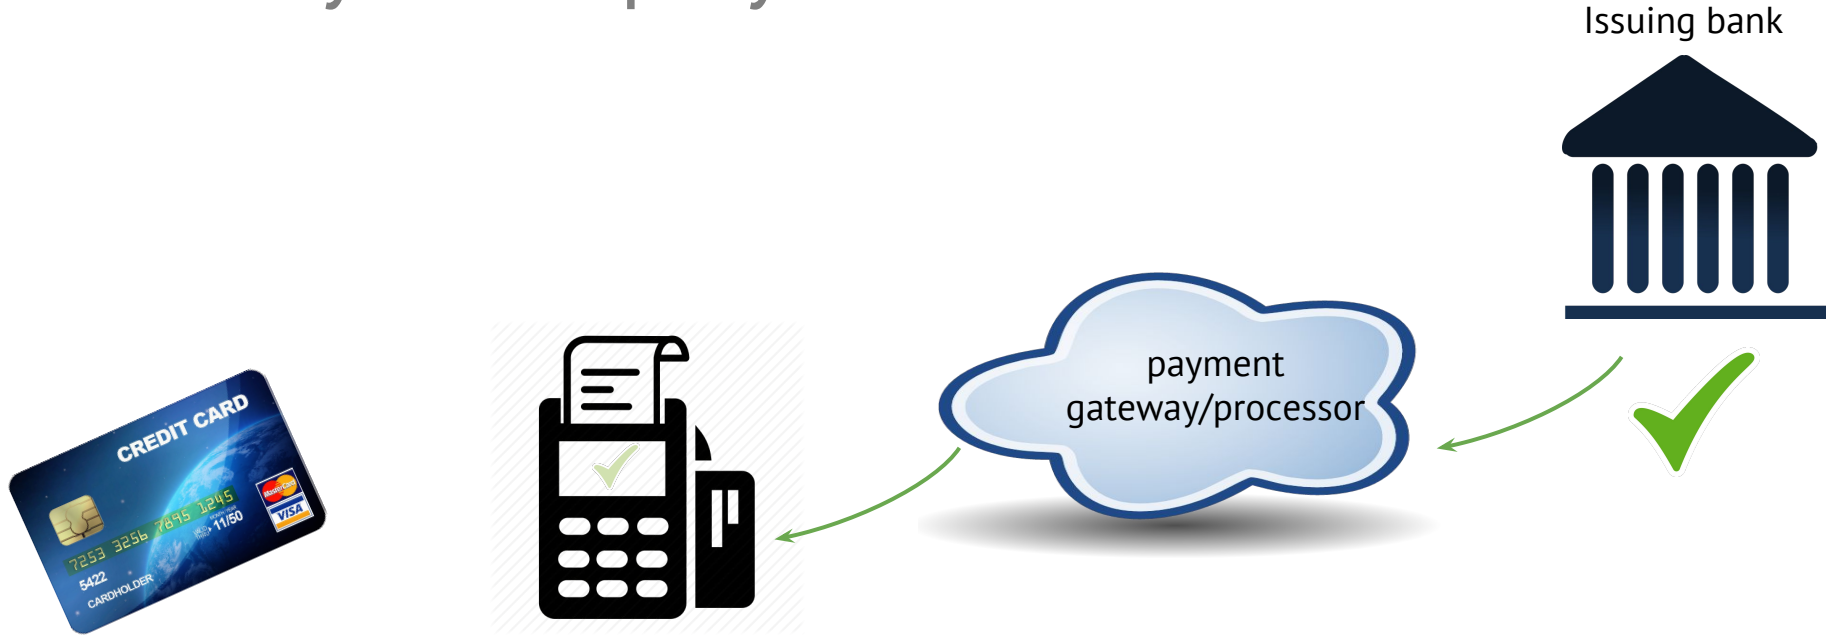
When you swipe your card



Issuing bank



When you swipe your card



Fraudster enters the play



Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play

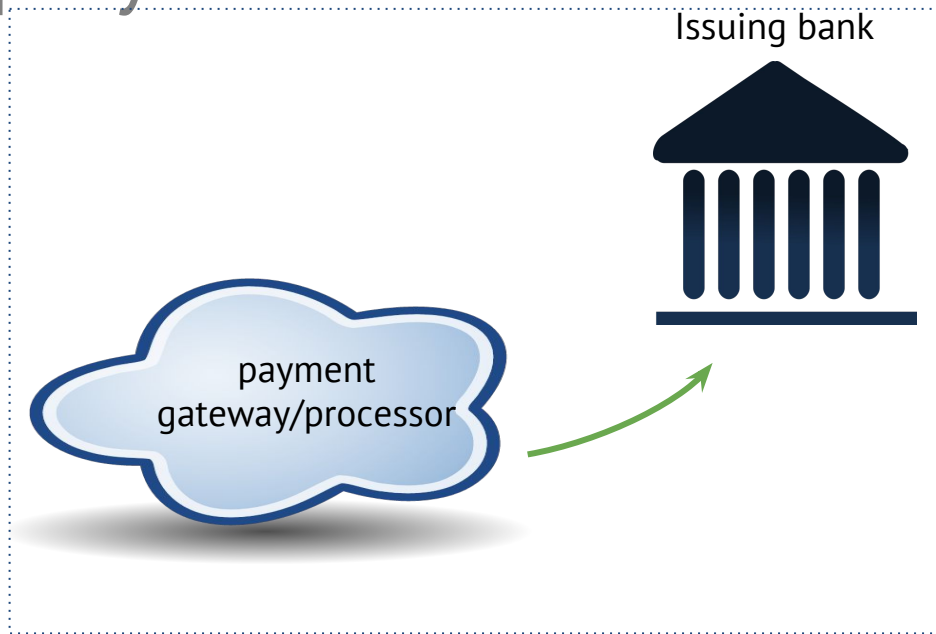


Issuing bank



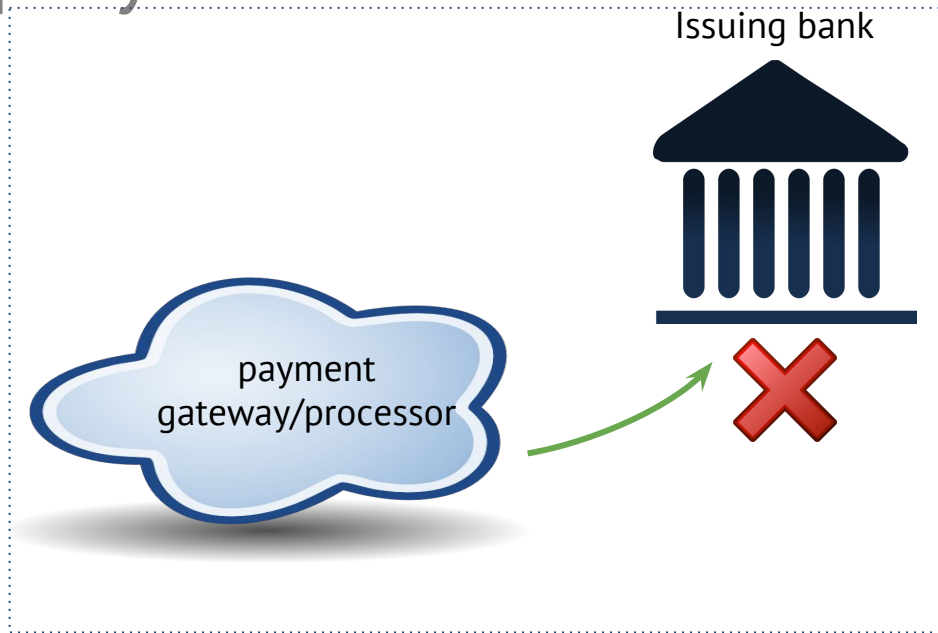
...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

Fraudster enters the play



...
7253 3256 7845 1245	12/17	055
...

But I protect my card!

“Where there is value, there is fraud”

data breaches
the fraud landscape

ATM machines



Well, at least one is an ATM machine...!

This is a card skimming device!




before



after

What happened? Mini card scanner...



..complete with camera and fake keyboard




If you avoid ATMs...

...you can still be targeted

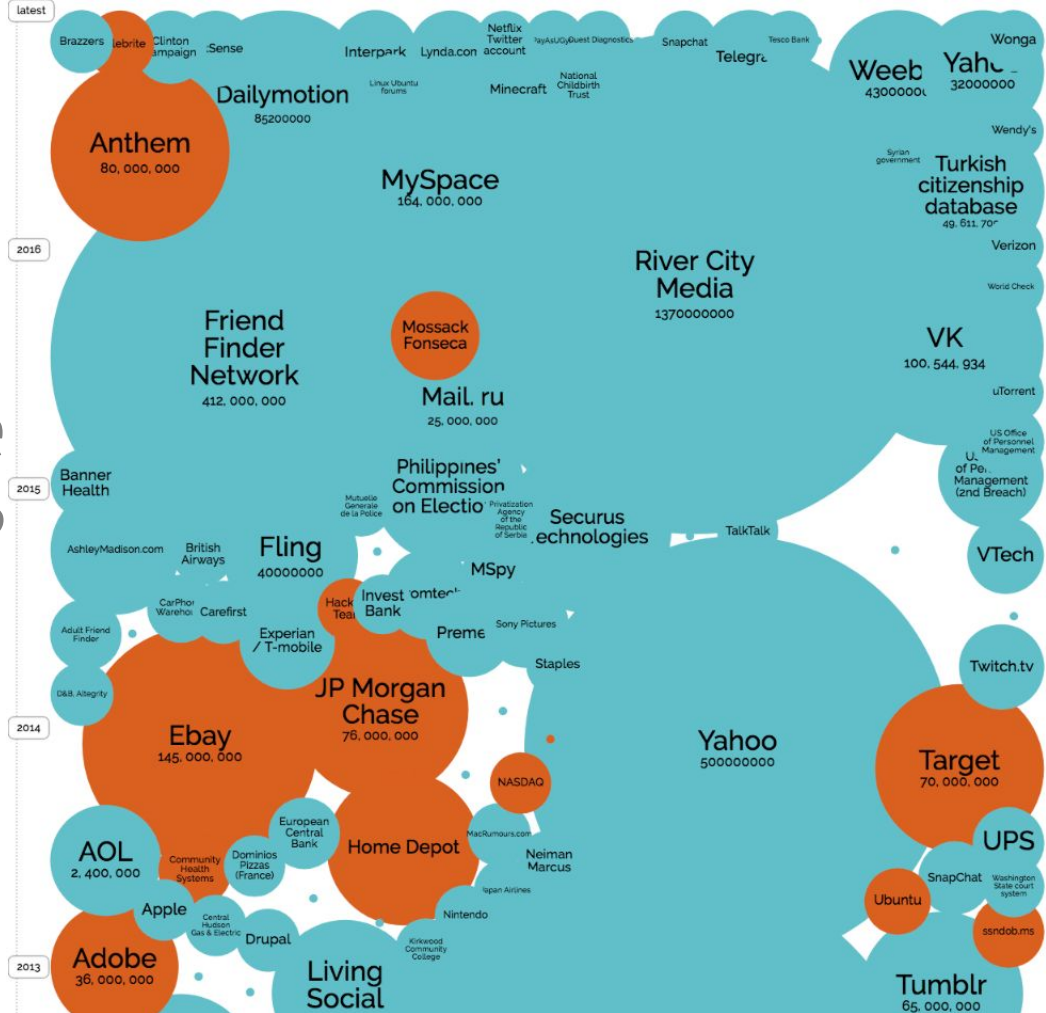
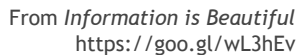


...at an astonishing speed



<https://www.youtube.com/watch?v=y83ZgzuFBSE>

How serious is the data breach landscape?



If my card is targeted, who is liable?



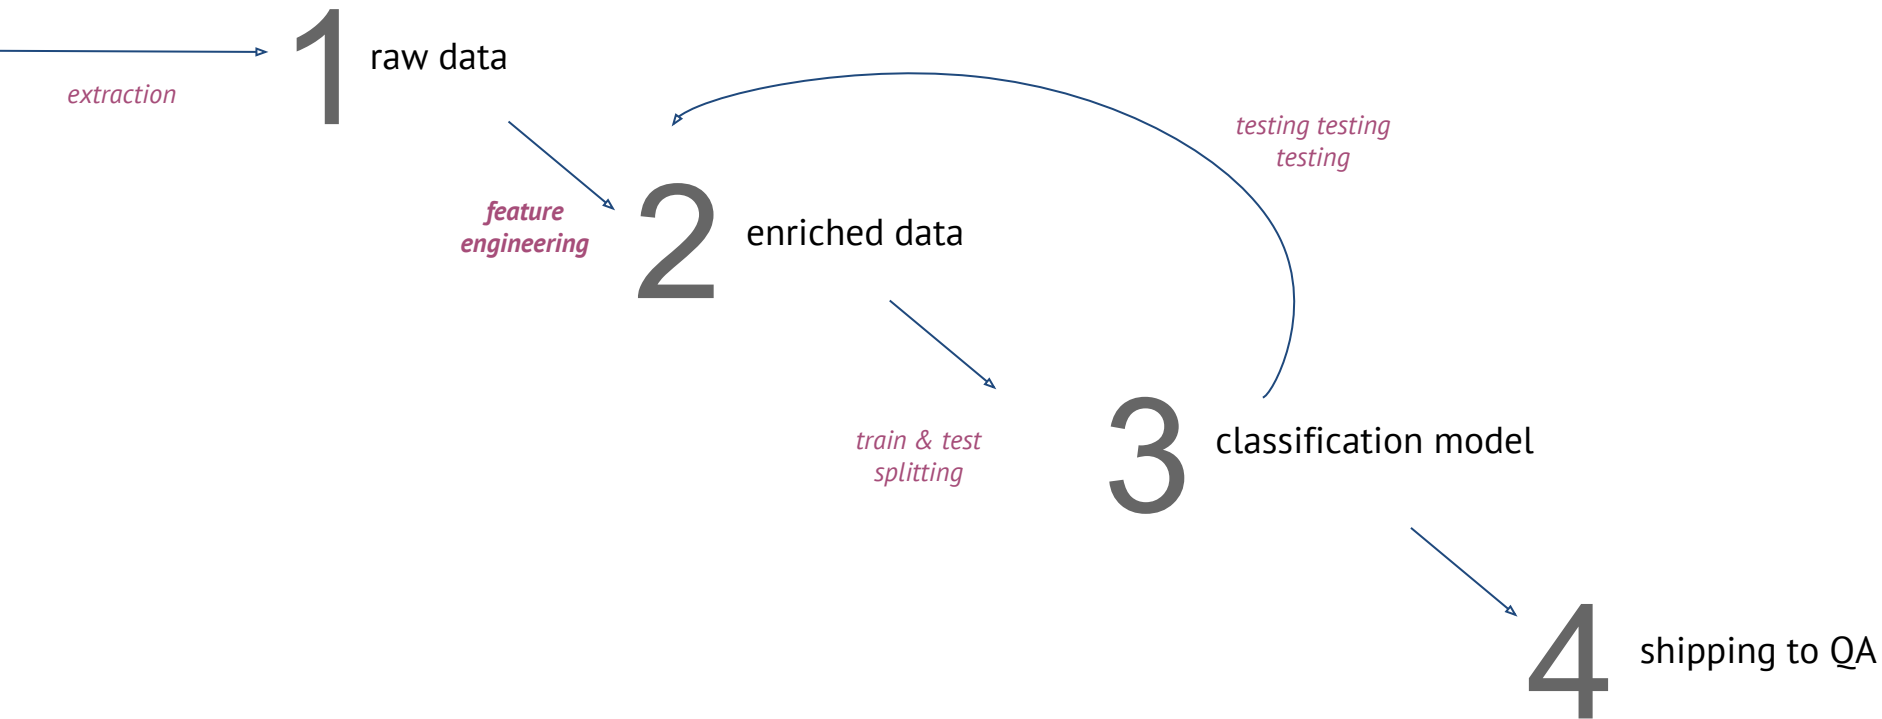
Machine learning insights

A - do you trust your labels?


B - how do you evaluate your model?

problems you won't usually encounter in kaggle competitions

A Data Science pipeline



A - When will you know it was fraud?



- You call your bank and request a chargeback
- You are reimbursed - happy ending!

Merchant is notified to pay the bill
30 days after transaction happened



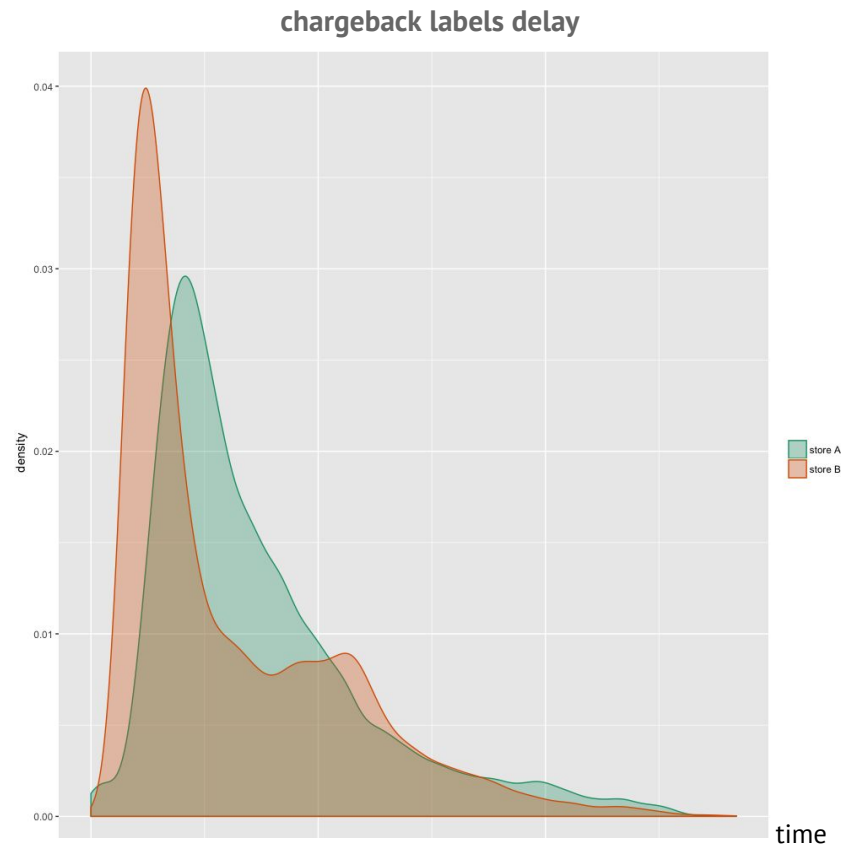
FACT: the classification model failed to identify a fraudulent transaction

PROBLEM: has the wrong label been fed into the model already?

A - Do you trust your labels?

How long are you willing to wait to get the purest labels?

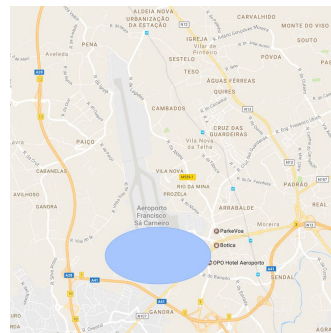
*The answer is a compromise between
more recent data and the correct labels.*



B - Tell me your metric

We deal with unbalanced datasets that range from

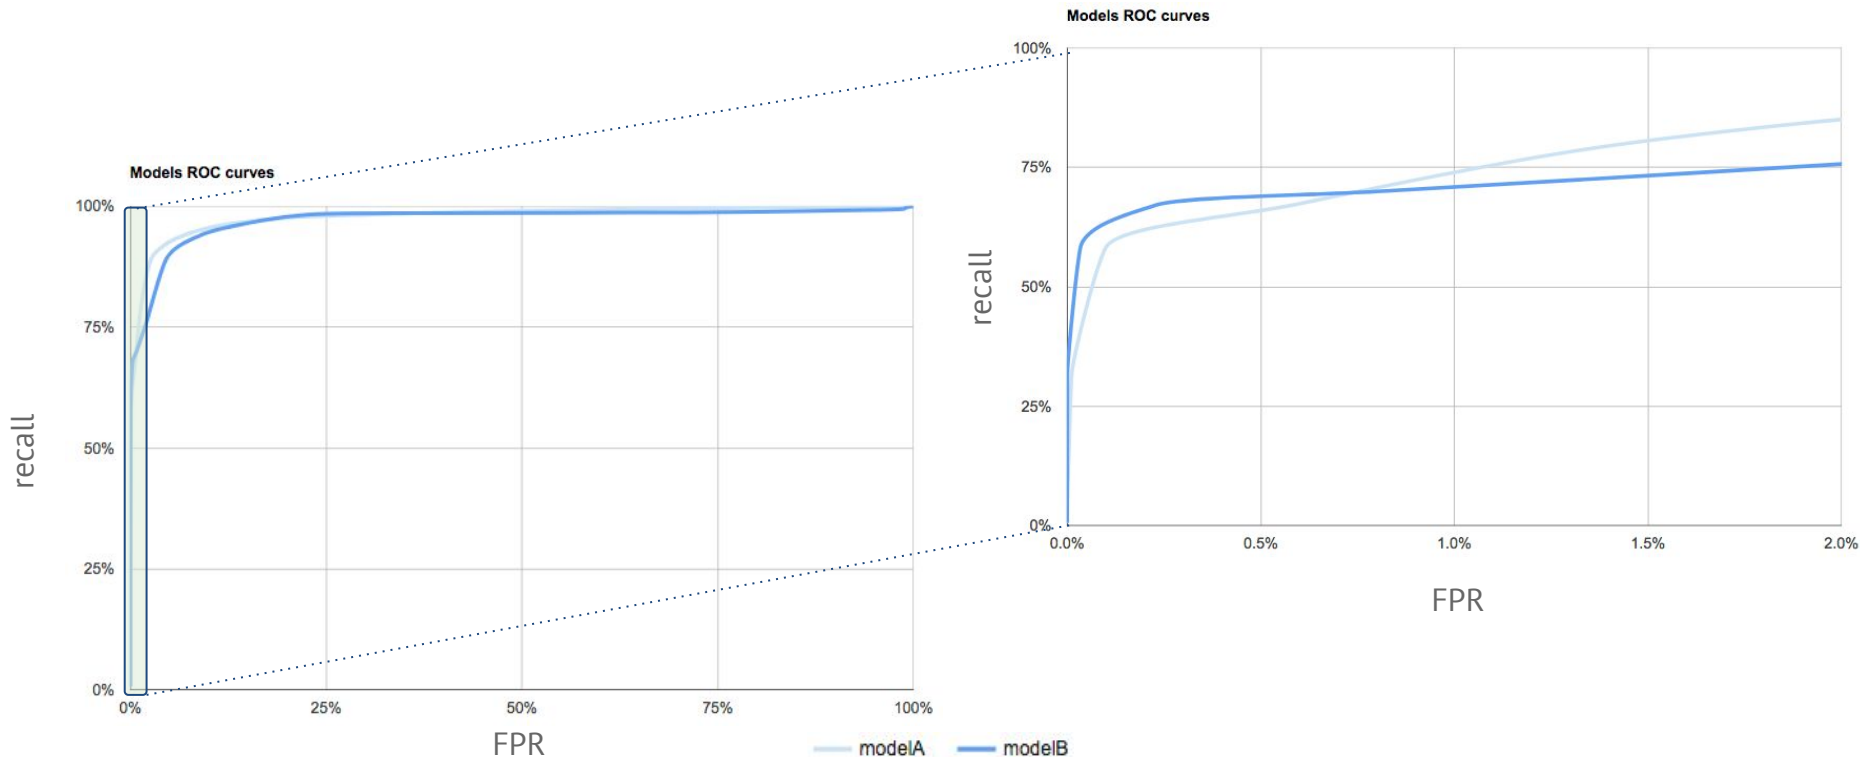
1:10 to **1:1000000**



A dummy classifier can reach over 90% in accuracy and miss all the fraud!

B - Decoding ROC curves

synthetic data




B - Scoring transactions in e-commerce

Though this is a classification problem, our models produce a score: from 0 to 1000.

Score 930 **probably fraud** send to review

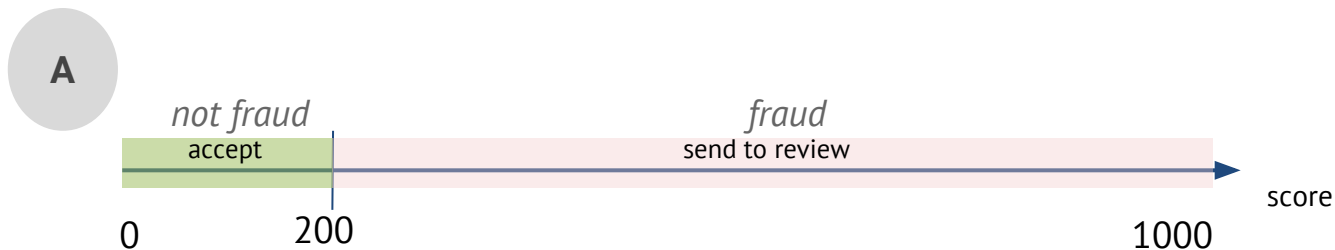
Score 20 **probably legit** approve

How do you define the threshold for **fraud** and **not fraud**?




B - How would you select a score threshold?

Though this is a classification problem, our models produce a score:



Catches most fraud
but also many
legitimate cases




Fails to catch some
fraud but should
trigger less FPs

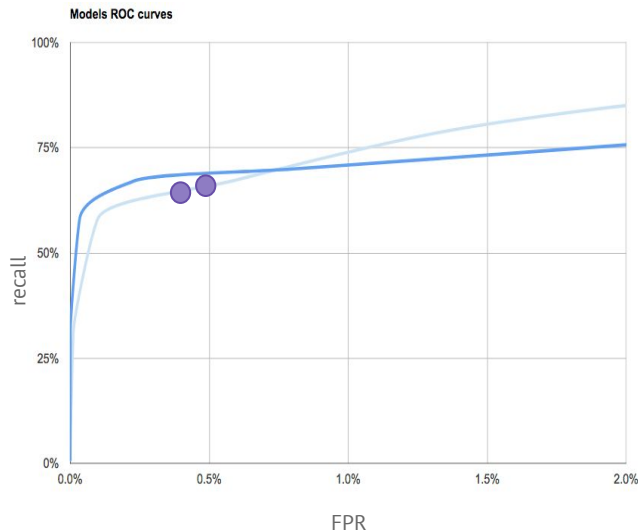

B - How would you select a score threshold?

But how do you decide between very close thresholds?

A



B



B - How do you evaluate your model?

Different thresholds will work for different SLAs.

Do you know what is the cost of a **FP** compared to a **TP**?

We can use **recall** to determine how much fraud we aim at catching and **FPR** to adjust how many FPs we are willing to allow for, cf.:

$$\text{recall} = \frac{TP}{TP + FN} \quad \text{and} \quad \text{FPR} = \frac{FP}{FP + TN}$$

Balance between catching fraud and keeping good customers satisfied with the service

B - How do you evaluate your model?

For those transactions which you are pretty certain are fraud, can you help the review process?



Do you still want to use **recall** to determine this threshold?

Are other metrics better suited for this? $\text{precision} = \frac{TP}{TP + FP}$

Machine Learning wrap-up

Summary

- Machine learning models face a spectacular adversarial problem in fighting fraud. Feature engineering and rapid deployments are key for success.
- In this talk we discussed two real-world problems:
 - in the business world, the **labels arrive late** and your models need to be deployed fast
 - **different metrics** for the algorithm performance are best suited **for different objectives**
- **Fighting fraud is very challenging but intellectually very stimulating!**



Thank-you!

Raquel H Ribeiro
Raquel.Ribeiro@feedzai.com

