

Deep Multimodal and Cross-modal Embeddings

DSPT #80 Webinar

David Semedo, Ph.D. df.semedo@campus.fct.unl.pt

Universidade NOVA de Lisboa, Portugal



NOVALINCS
LABORATORY FOR COMPUTER
SCIENCE AND INFORMATICS



FACULDADE DE
CIÊNCIAS E TECNOLOGIA
UNIVERSIDADE NOVA DE LISBOA

About me and my research

- Researcher at the **Web and Media Search** group, from NOVA LINCS since 2015.
- Ph.D. in **Deep Learning for Multimedia Understanding**.
- Topic:
 - “*Bridging Vision and Language over Time with Neural Cross-modal Embeddings*”.
- **Main interests:** multimodal machine learning, at the intersection of CV and NLP, neural networks and data mining.




Webinar outline

- Motivation
- Learning deep multimodal and cross-modal embeddings
- State-of-the-art and Applications
- Final remarks

The world is multimodal!

Vision and Language

Visual Commonsense Reasoning: Answer a question about an image and provide a rationale justifying the answer.



Why is [person1] pointing a gun at [person2]?

- a) [person1] wants to kill [person2]. (1%)
- b) [person1] and [person3] are robbing the bank and [person2] is the bank manager. (71%)**
- c) [person2] has done something to upset [person1]. (18%)
- d) Because [person2] is [person1]'s daughter. [person1] wants to protect [person2]. (8%)


b) is right because...

- a) [person1] is chasing [person1] and [person3] because they just robbed a bank. (33%)
- b) Robbers will sometimes hold their gun in the air to get everyone's attention. (5%)
- c) The vault in the background is similar to a bank vault. [person3] is waiting by the vault for someone to open it. (49%)**
- d) A room with barred windows and a counter usually resembles a bank. (11%)

The world is multimodal!

Different diagnostic exams

- Positron Emission Tomography (**PET**)
- Computer Tomography (**CT**)
- Magnetic Resonance Image (**MRI**)

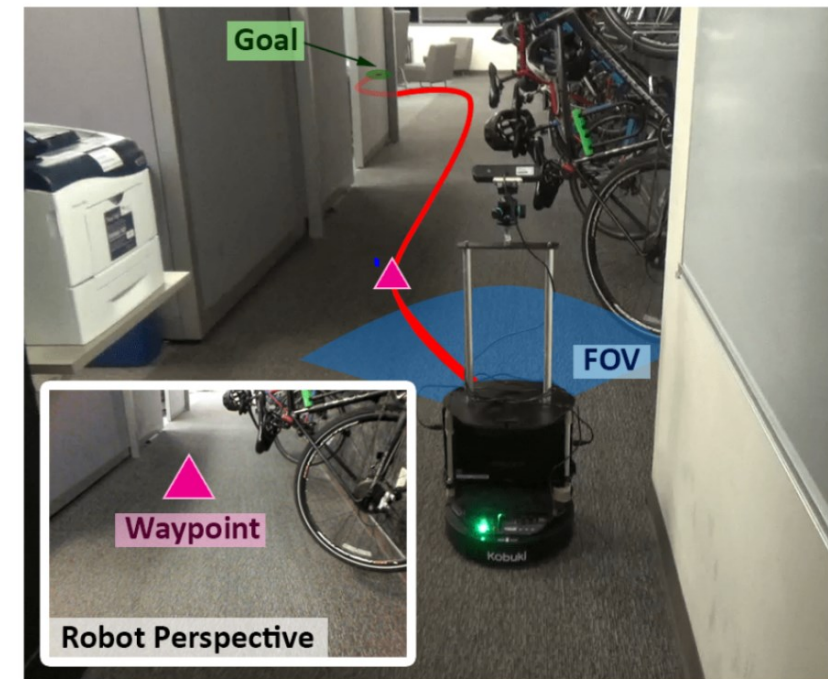


Tumor segmentation

The world is multimodal!

Autonomous Navigation

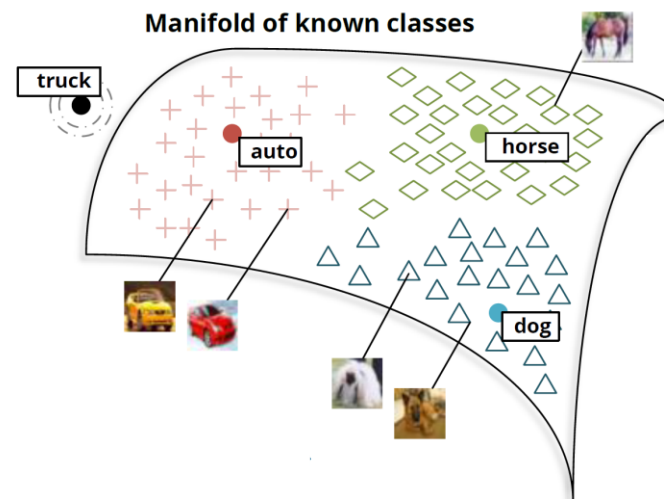
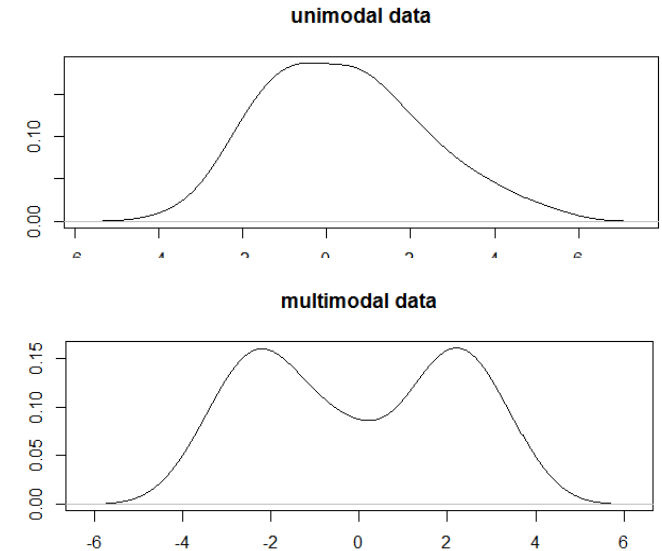
- Vision (e.g. RGB, Infrared Cameras)
- Sensors (e.g. proximity, depth)
- Speech
- ...



Motivation - Definitions

Multimodal: “Having or involving several modes, modalities, or maxima”

Multimodal Embedding: Vector representation of a given document, in a multimodal space



Multiple modalities in ML pipelines

Why would we want to bridge two (or more) modalities?

- ✓ Combining information from multiple information sources!



Description:

“Deep Time light show at
Edinburgh Festival, 2016”

Structured vs. Unstructured data

Tabular data – After vectorizing, each dimension potentially discriminates different samples.

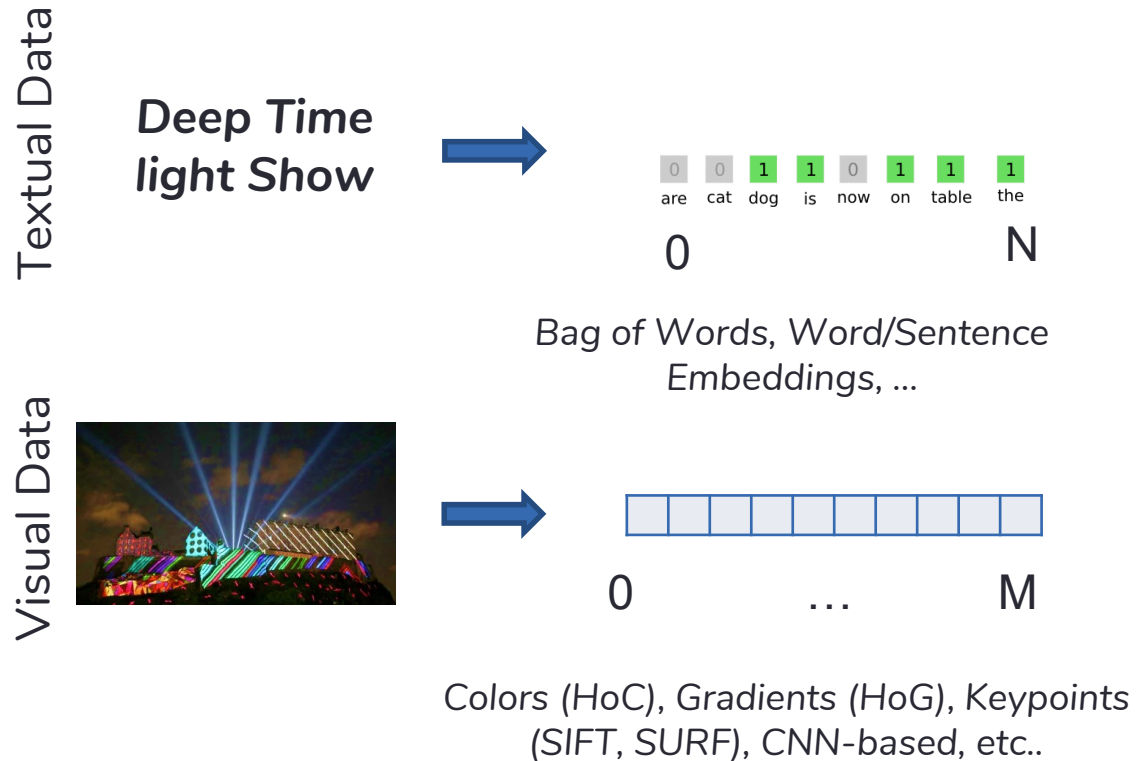
	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

Images (RGB - 3-channel width x height matrix)



Why not just concatenate representations?

Heterogeneous Representations




Early fusion (e.g. concatenation)



- Representations need to be well-aligned;
- Single model for both modalities;
- Limited capability to unveil complex correlations between modalities.

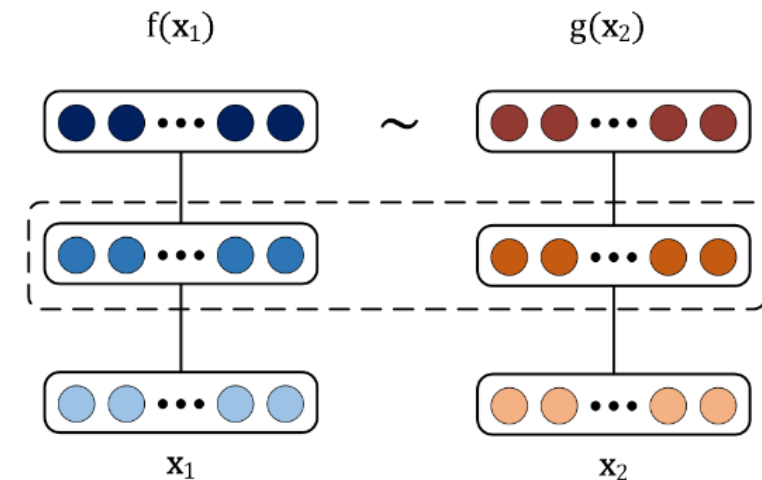
Multimodal Embeddings – Joint Representations

- **Merges information** from all input modalities in a single embedding after non-linear transformations (**Late fusion**);
- Multimodal **data present at training and inference time.**




Multimodal Embeddings – Coordinated Representations

- Learn **two separate spaces**, that are constrained to be aligned:
 - Minimize distance metric, maximize correlation, etc..
- **Independent projection functions** - copes with missing modalities!



Linear Cross-modal Embedding

Given a projection basis A for texts and B for images, project both modalities to correlated subspaces.



Canonical Correlation Analysis

CCA objective (maximize):

$$\rho = \frac{E[uv]}{\sigma_u \sigma_v} = \frac{E[uv]}{\sqrt{E[u^2]E[v^2]}}$$

Going beyond linear transformations

Do we know *a priori* *how* any two modalities are correlated?

We may have a hint, but not exactly
how to mathematically express that hint.


State-of-the-art:

- ✓ Use neural networks to learn **non-linear projection functions.**

Universal Approximation Theorem:

We just need enough width (neurons) or depth (hidden layers)!

Multimodal Autoencoder



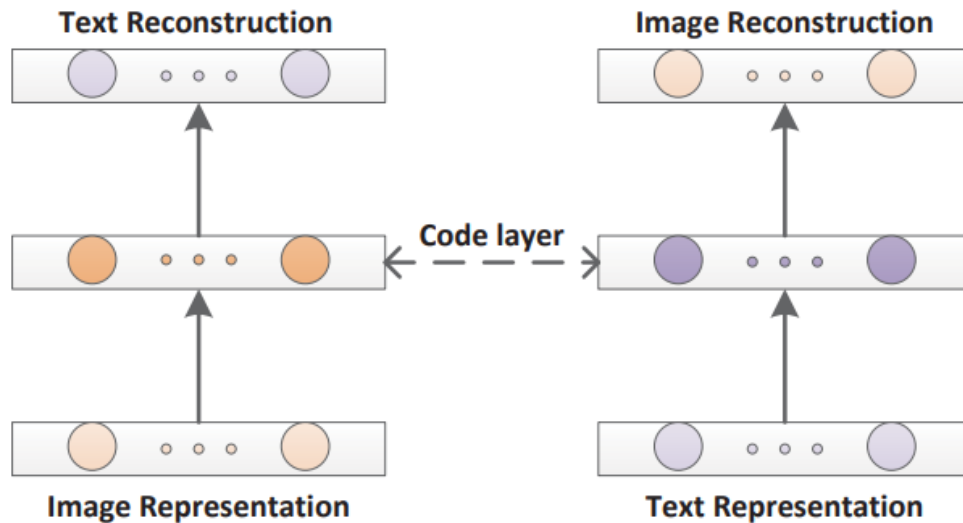
Loss: l_2 reconstruction error.

What can we do with such architecture/embedding?

- Extend it to more than 2 modalities;
- Multimodal queries;
- Use as features on a classifier.

Not a cross-modal embedding yet

Correspondence Autoencoder



Loss: l_2 reconstruction error.

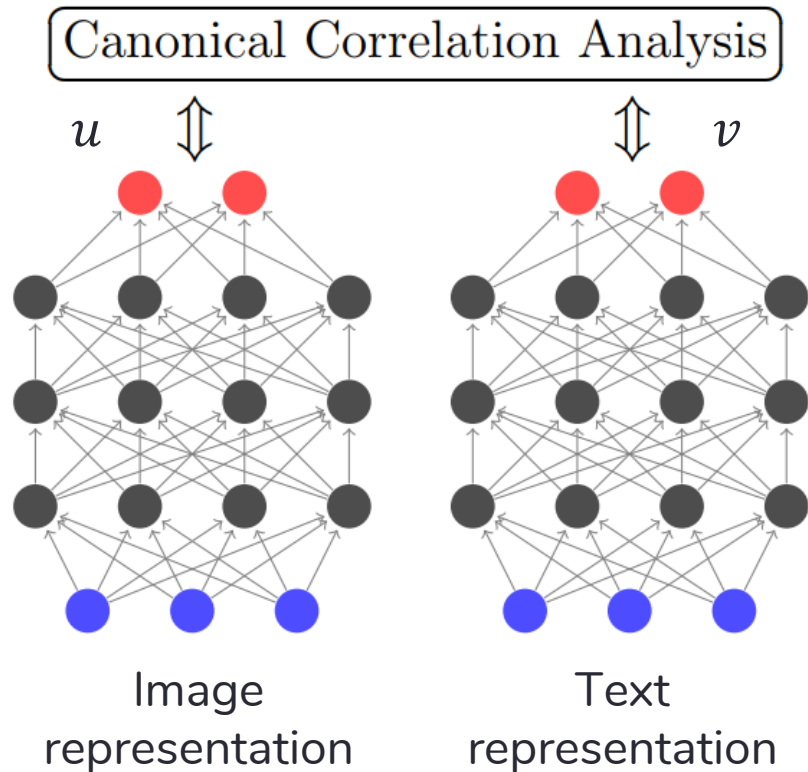
Same input reconstruction principle as in an Autoencoder, but ...

Weights of the Intermediate layer are shared.

This means that each modality network can be used independently!

Deep Canonical Correlation Analysis

Neural network projection function, with the CCA objective as loss function.



Learns a representation by **maximizing correlation**.

$$\rho = \frac{E[uv]}{\sigma_u \sigma_v} = \frac{E[uv]}{\sqrt{E[u^2]E[v^2]}}$$

Unsupervised approach, but it's quite an effective architecture!

Going beyond pairwise-correlations

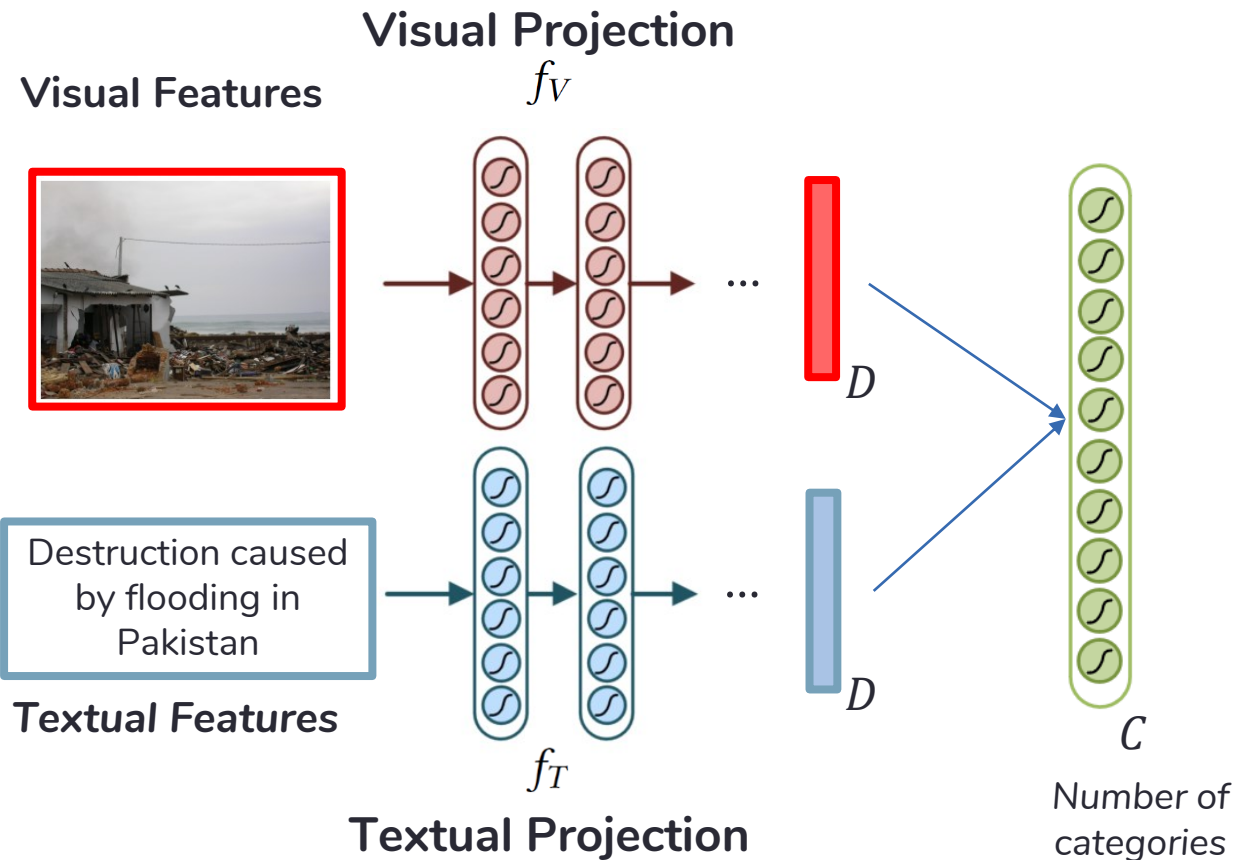
State-of-the-art:

- ✓ Use neural networks to learn **non-linear projection functions**.
- ✓ **Supervised**: Consider category/class information to help structuring data.

Same class



Supervised Cross-modal Embeddings – Classifier as proxy



Multi-class setting:
softmax layer

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Loss: **categorical cross-entropy**


Multi-label setting:
logits layer (multiple regression functions)

Loss: **binary cross-entropy**

$$[y_n \cdot \log \sigma(x_n) + (1 - y_n) \cdot \log(1 - \sigma(x_n))]$$

sigmoid

Supervised Cross-modal Embeddings – Metric Learning

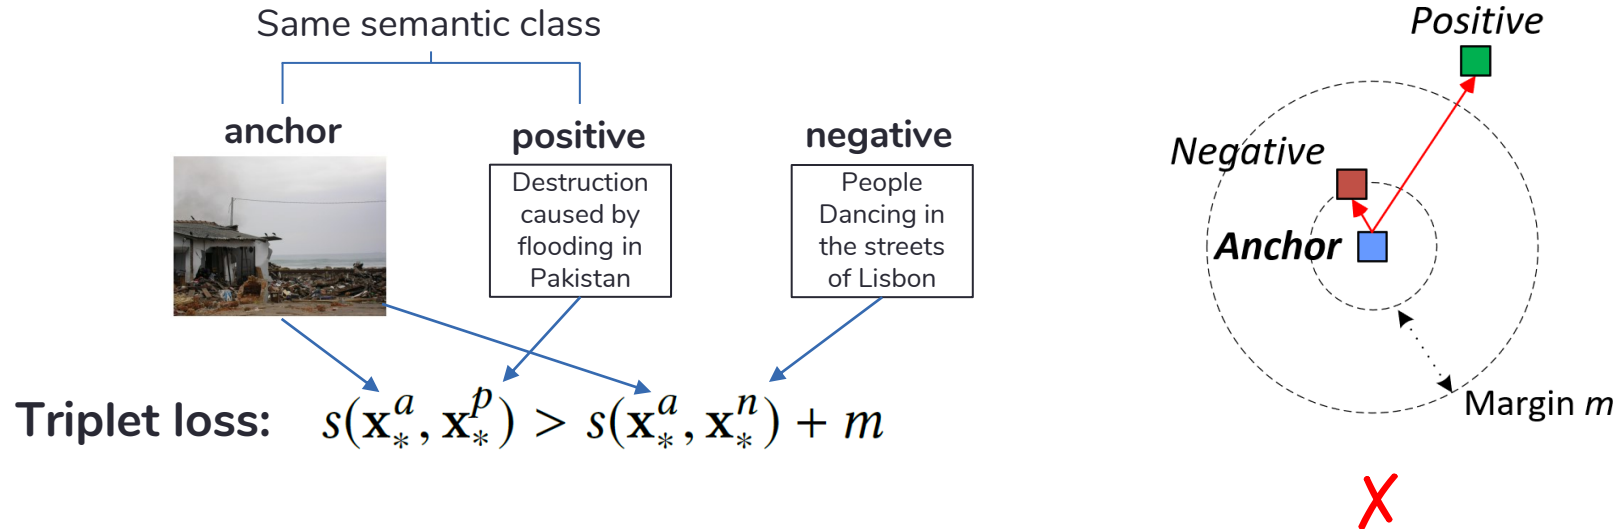


Loss: **ranking loss**

The loss directly imposes the structure we want on the embedding space.

Example of a Ranking loss - Triplet Loss

Directly enforces the desired embedding structure.




Example of a Ranking loss - Triplet Loss

$$s(\mathbf{x}_*^a, \mathbf{x}_*^p) > \underbrace{s(\mathbf{x}_*^a, \mathbf{x}_*^n)} + \boxed{m}$$

Constraint is formulated as
a **differentiable function**:

$$loss(\mathbf{x}_*^a, \mathbf{x}_*^p, \mathbf{x}_*^n) = \max(0, m - s(\mathbf{x}_*^a, \mathbf{x}_*^p) + s(\mathbf{x}_*^a, \mathbf{x}_*^n))$$

Then, **sample several triplets** and keep enforcing the constraint over those triplets:



$$loss = \underbrace{loss(\mathbf{x}_*^a, \mathbf{x}_*^p, \mathbf{x}_*^n)}_{\text{Triplet 1}} + \dots + \underbrace{loss(\mathbf{x}_*^a, \mathbf{x}_*^p, \mathbf{x}_*^n)}_{\text{Triplet N}}$$

Adaptive Triplet Loss: scheduled optimization approach – increases expressiveness (ref. below).

Some results on cross-modal retrieval

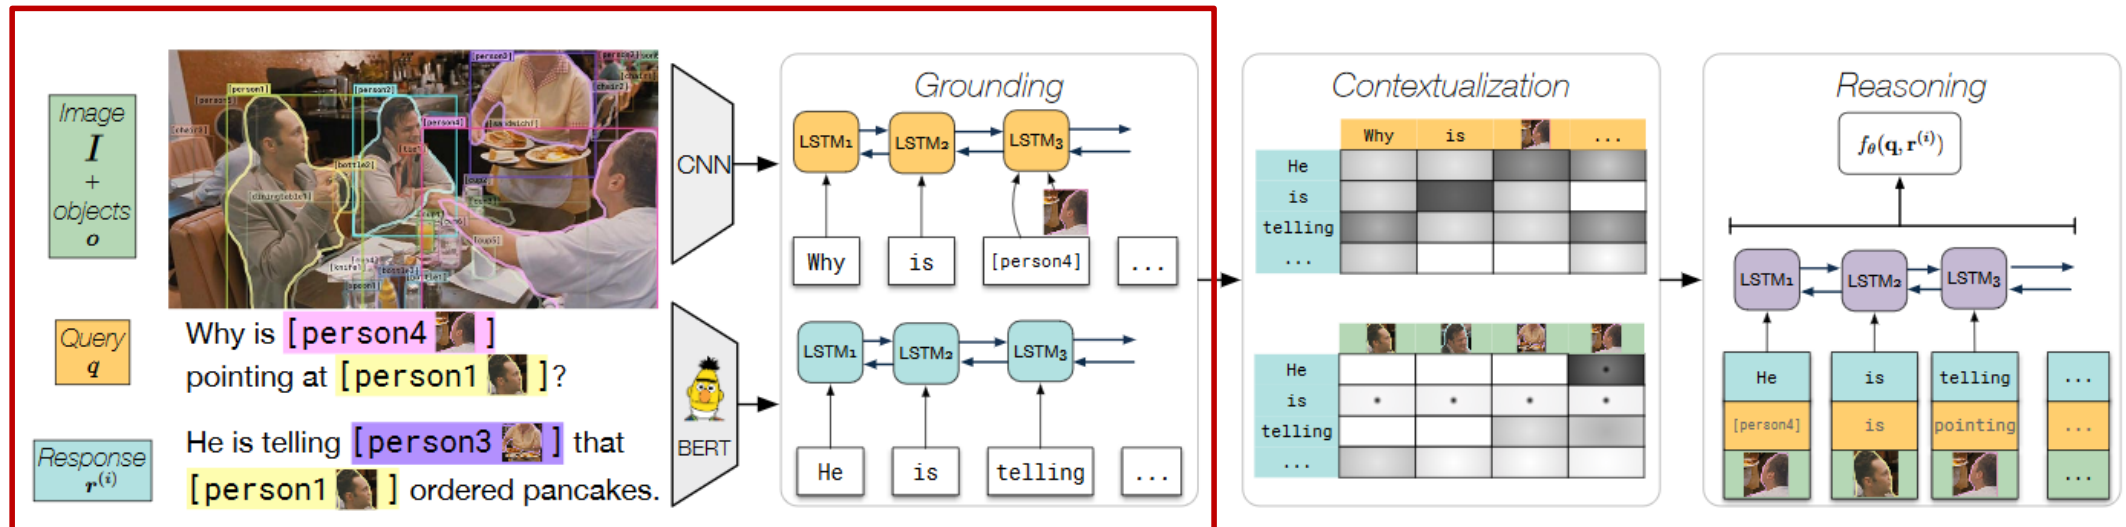
Results (mean Average Precision) on 3 benchmark datasets.

Method		Pascal Sentences			NUS-WIDE-10k			Wikipedia		
		$I \mapsto T$	$T \mapsto I$	Avg	$I \mapsto T$	$T \mapsto I$	Avg	$I \mapsto T$	$T \mapsto I$	Avg
Linear	<u>CCA</u>	0.203	0.208	0.206	0.167	0.181	0.174	0.298	0.273	0.286
	CFA	0.476	0.470	0.473	0.406	0.435	0.421	0.319	0.316	0.318
	KCCA	0.488	0.446	0.467	0.351	0.356	0.354	0.438	0.389	0.414
	LGCFL	0.539	0.503	0.521	0.453	0.485	0.469	0.466	0.431	0.449
	JRL	0.563	0.505	0.534	0.466	0.499	0.483	0.479	0.428	0.454
Correspondence Autoencoder	<u>Corr-AE</u>	0.532	0.521	0.527	0.441	0.494	0.468	0.442	0.429	0.436
Deep CCA	<u>DCCA</u>	0.568	0.509	0.539	0.452	0.465	0.459	0.445	0.399	0.422
	CMDN	0.544	0.526	0.535	0.492	<u>0.542</u>	0.517	0.487	0.427	0.457
Classifier-based	<u>Deep-SM</u>	0.560	0.539	0.550	0.497	0.478	0.488	0.478	0.422	0.450
	ACMR	0.538	0.544	0.541	<u>0.519</u>	<u>0.542</u>	<u>0.531</u>	0.468	0.412	0.440
	CCL	<u>0.576</u>	<u>0.561</u>	<u>0.569</u>	0.481	0.520	0.501	<u>0.505</u>	<u>0.457</u>	<u>0.481</u>
Adaptive Triplet loss	<u>SAM</u>	0.637	0.643	0.640	0.563	0.594	0.579	0.518	0.457	0.487

State-of-the-art and Applications

Revisiting Visual Commonsense Reasoning: Answer a question about an image and provide a rationale justifying the answer.

Model Architecture:

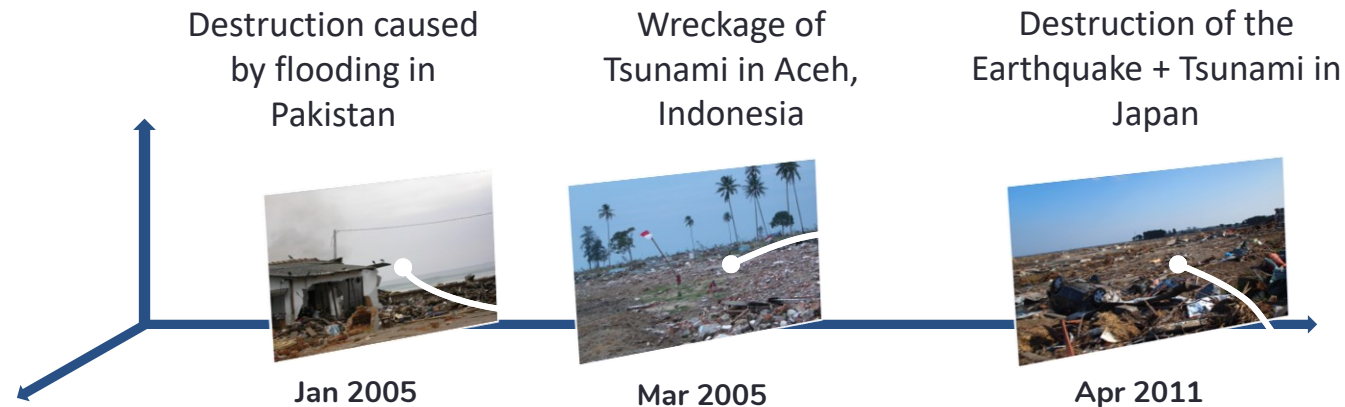


Bringing extra information to these embeddings

Adding the time dimension:

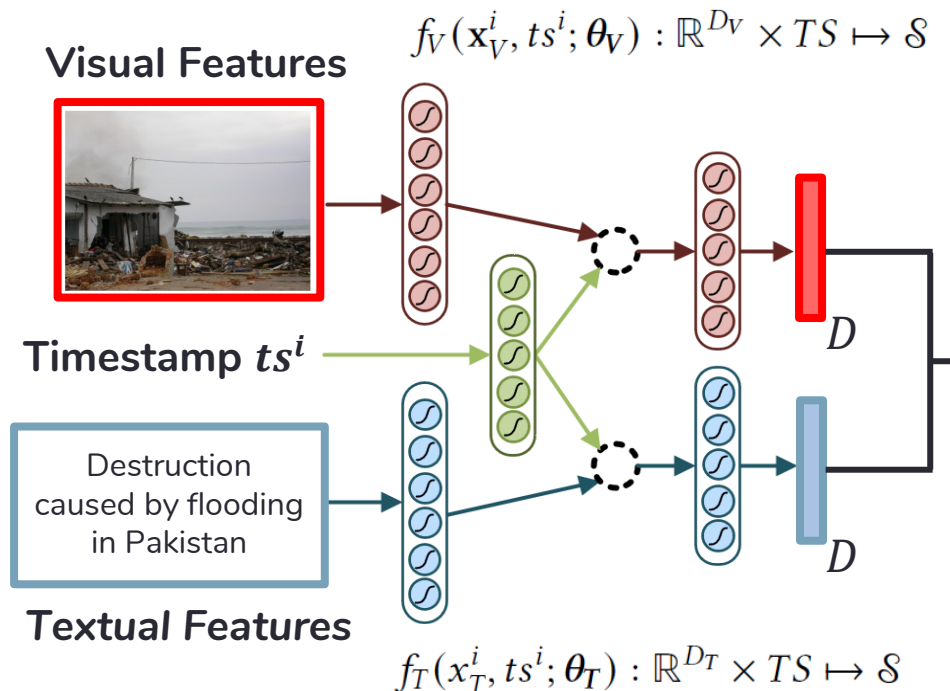
- Modeling the evolution of vision and language interactions.

Diachronical cross-modal embedding



Diachronic Cross-modal Embeddings

Architecture



- ✓ **Conditions** representations on **time** (continuous input);
- ✓ **Jointly learned**.

$$\mathcal{L}(x_*^a, x_*^p, x_*^n; \theta) = \mathcal{L}_{inter}(x_*^a, x_*^n; \theta) + \mathcal{L}_{intra}(x_*^a, x_*^p; \theta)$$

Triplet loss


Temporal Alignment
to preserve data
original timeline

Analysis of Semantic Dispersion over Time on **20 years** of data

Tsunami
Indonesia



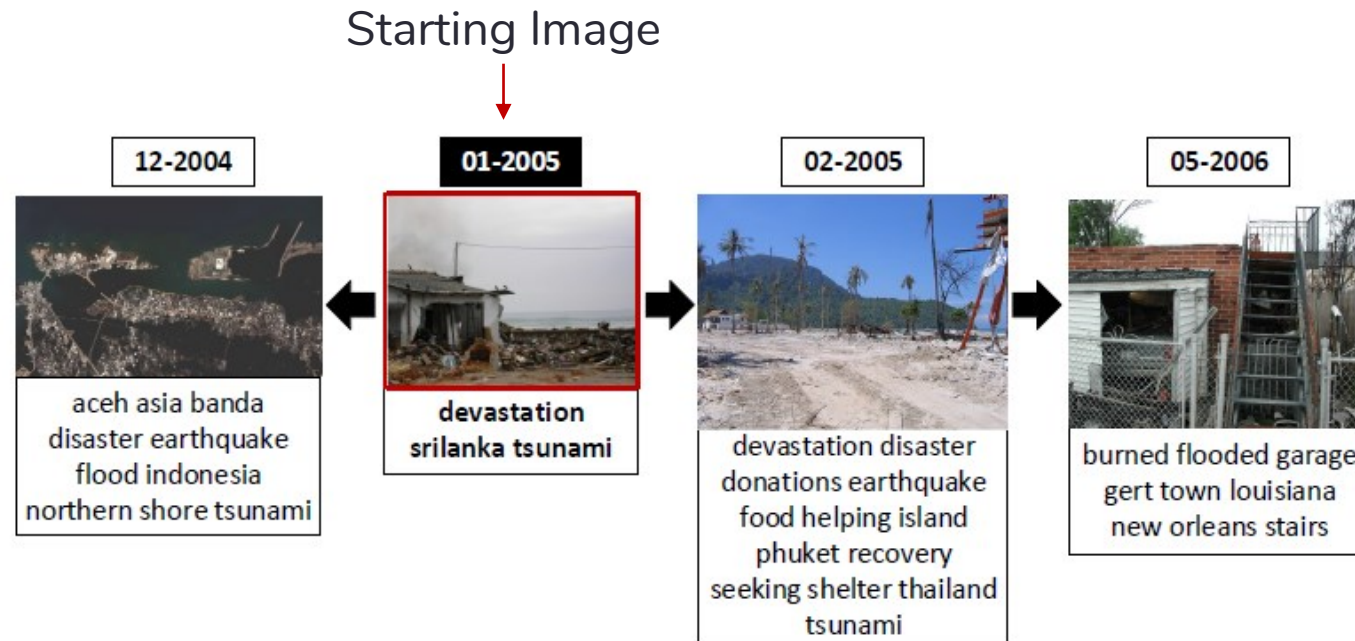
Tsunami
Japan



Corresponds to the timestamps
of the images

Cross-modal Evolution – Summarizing Trajectories


Automatically generate summaries given a single Image/Text,
based on temporal trajectories



State-of-the-art and Applications

12-in-1: Multi-Task Vision and Language Representation Learning:

- ✓ Based on **ViLBERT** - BERT extension to Jointly Represent Images and Text.



Visual Question Answering

What color is the child's outfit? **Orange**

Referring Expressions

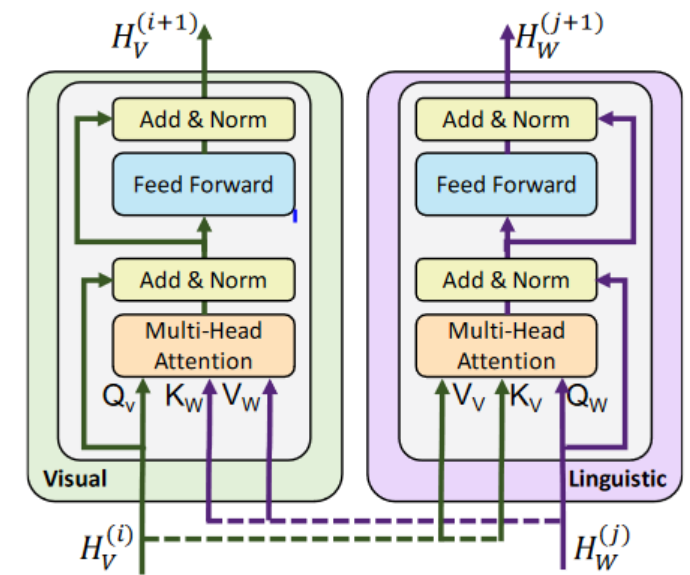
child sheep basket people sitting on chair

Multi-modal Verification

The child is petting a dog. **false**

Caption-based Image Retrieval

A child in orange clothes plays with sheep.



Co-attention transformer layer

Information exchange between modalities.

Research direction: Multi-task learning + several datasets.

Summary and Key Takeaways

Multimodal and Cross-modal embeddings can effectively combine information from multiple modalities.

- We've covered the basic building blocks to learn effective multimodal embeddings.
- Feed ML models to address down-stream tasks.

Further reading and pointers

- Awesome multimodal ML list ([Link](#));
- Multimodal Machine Learning tutorial ([Link](#));
- Nice survey: Baltrušaitis, Tadas et al. **Multimodal Machine Learning: A Survey and Taxonomy**, PAMI 2019;
- Some relevant conferences: ACM MM, CVPR, NeurIPS, ICCV, ICML, ...
- Deep Learning book, Ian Goodfellow, Yoshua Bengio and Aaron Courville ([Link](#)).

Thank you!

David Semedo, Ph.D. df.semedo@campus.fct.unl.pt

Universidade NOVA de Lisboa, Portugal

