

A Multi-Cloud Infrastructure with Kubernetes, Prometheus and other CNCF Projects

Ricardo Rocha
Computing Engineer, CERN

Data Science Portugal

Ricardo Rocha

Computing Engineer, CERN Cloud Team

Kubernetes and Containers, Networking and SDN

GPUs and other Accelerators, Machine Learning

Cloud Native Computing Foundation (CNCF)

Representative of CERN in the CNCF and End User Community

Member of the CNCF Technical Oversight Committee (TOC)

<https://www.cncf.io/people/technical-oversight-committee/>

Lead of the CNCF Research User Group

<https://github.com/cncf/research-user-group>





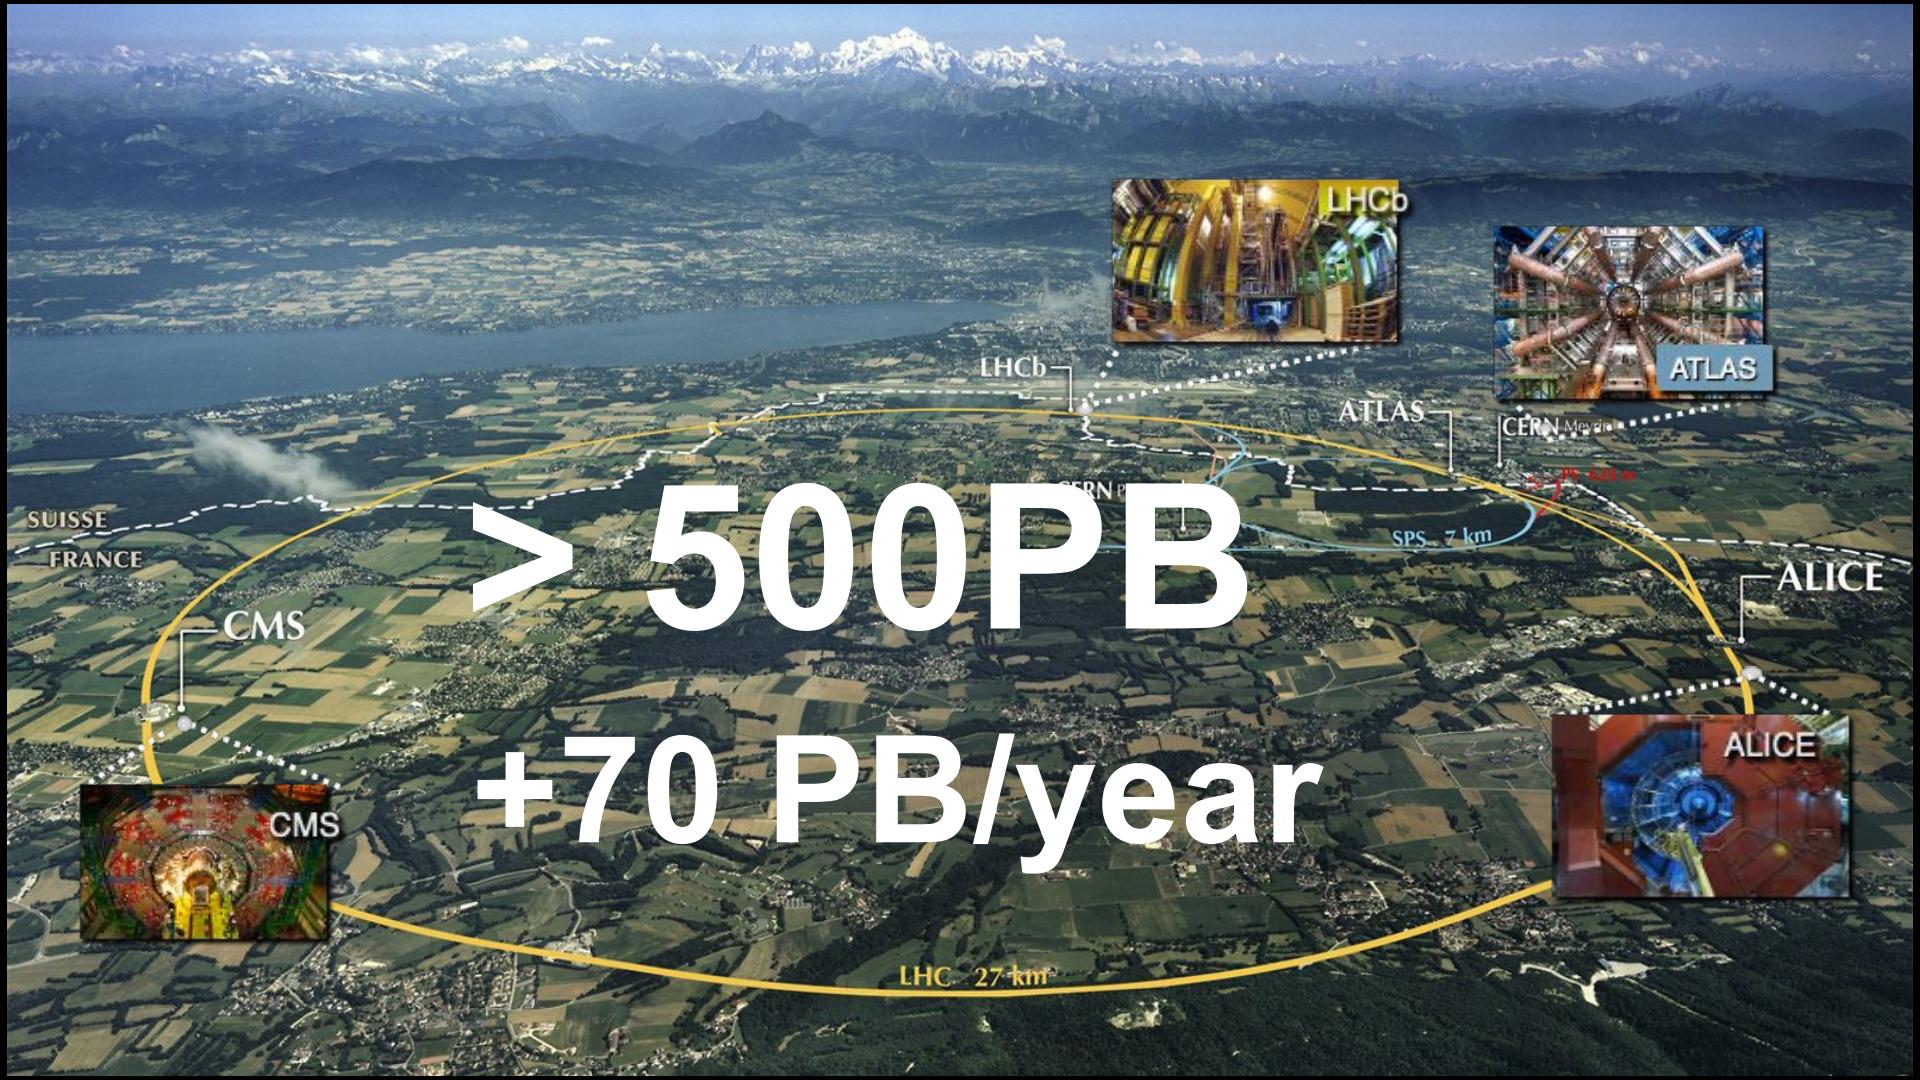
Founded in 1954

Fundamental Science

What is 96% of the universe made of?

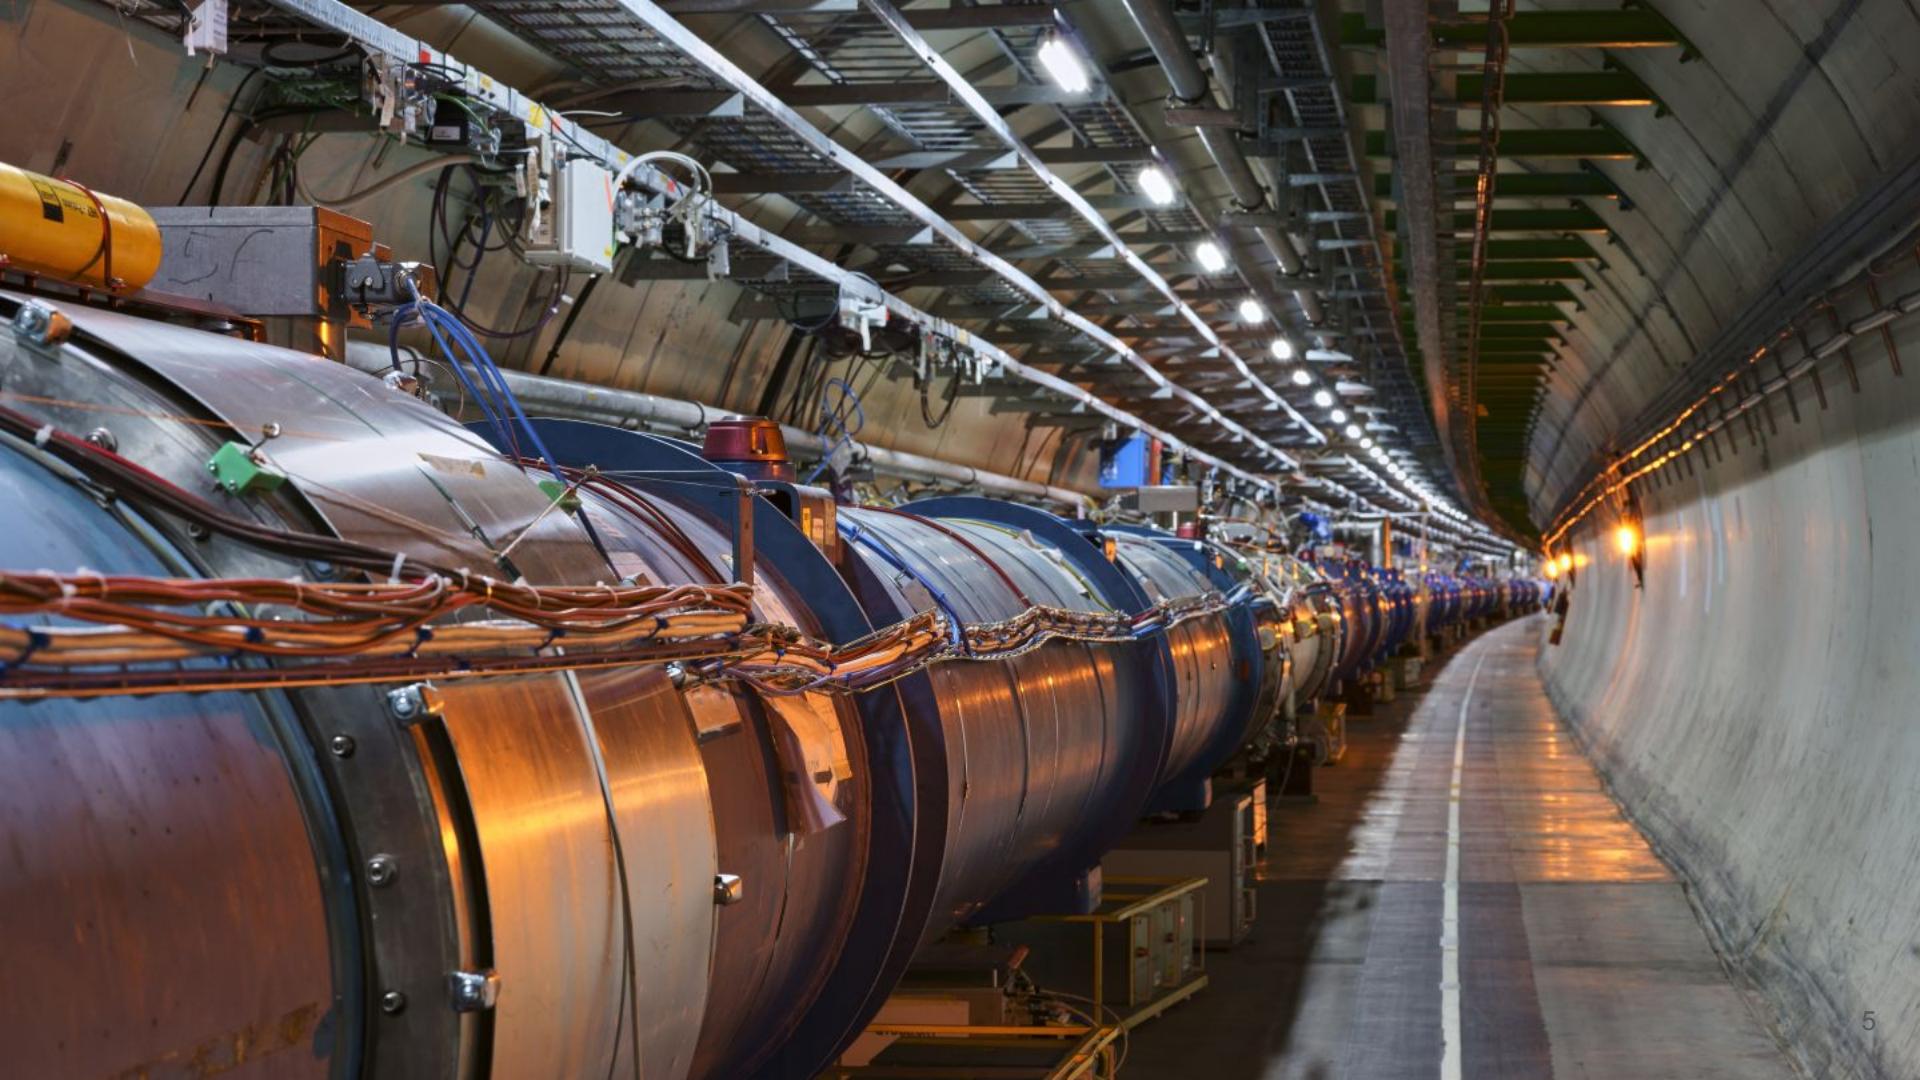
What was the state of matter just after the Big Bang?

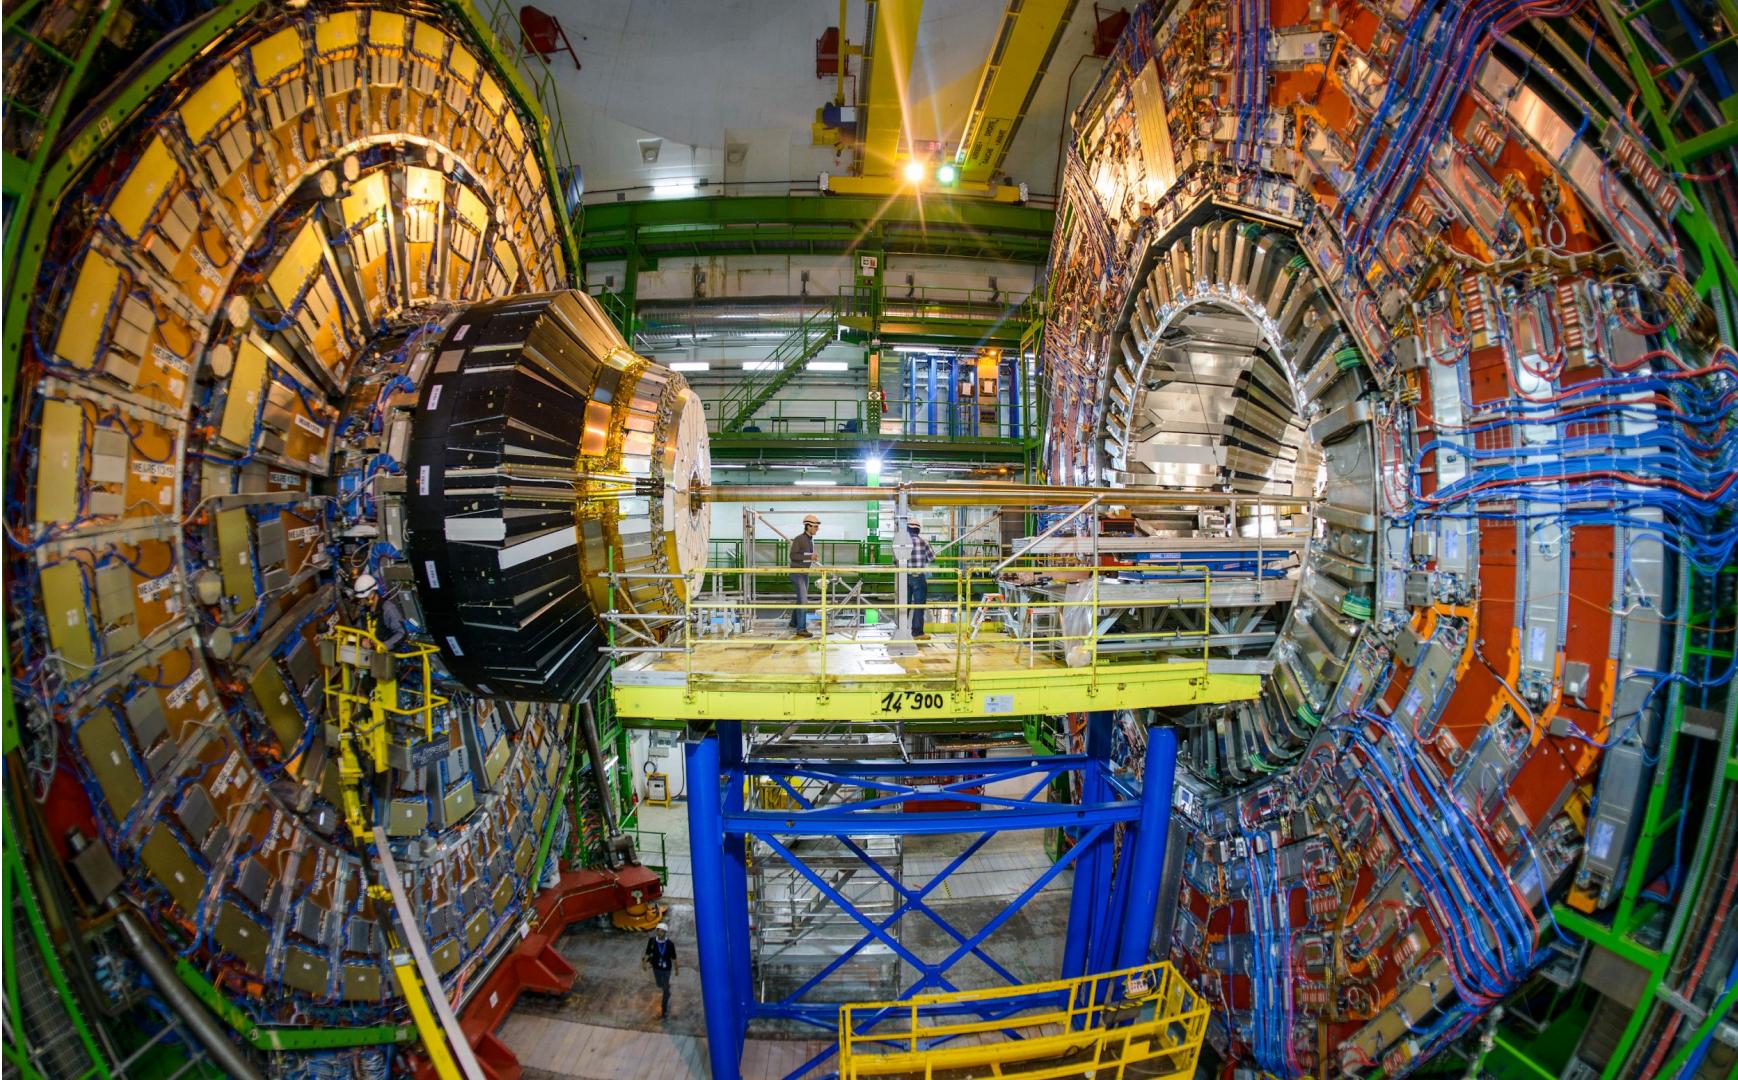
Why isn't there anti-matter in the universe?



**> 500 PB
+70 PB/year**







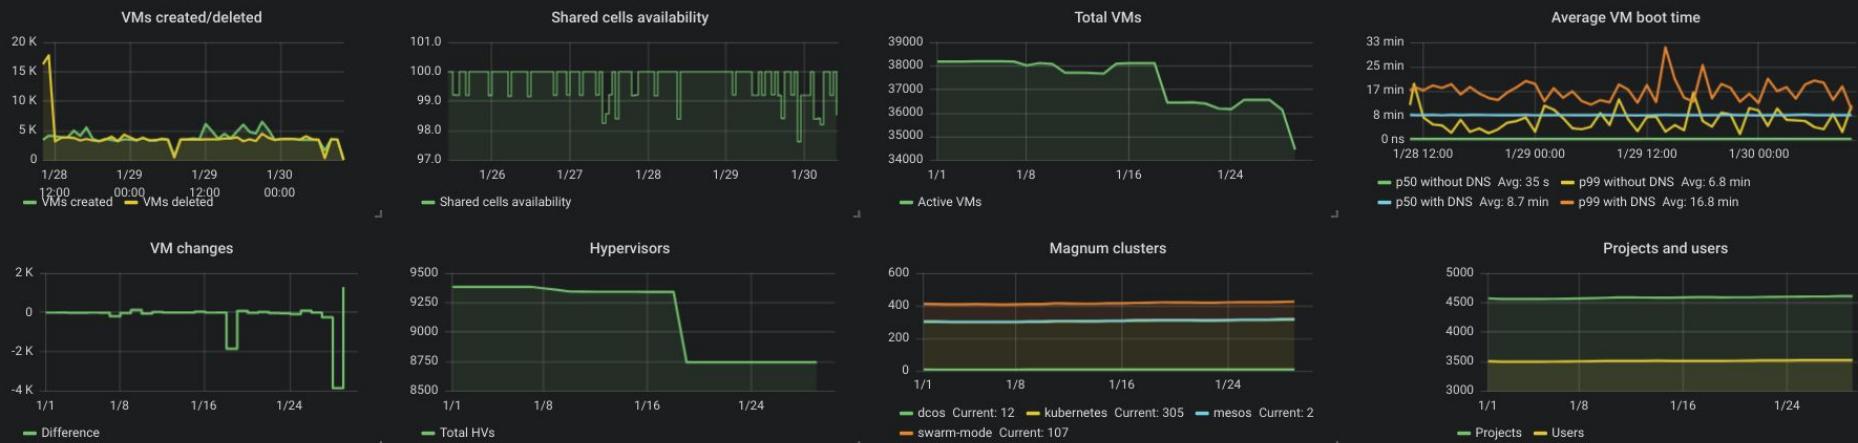
Cloud resources



Openstack services stats



Resource overview by time



CERN Kubernetes

Statistics on clusters running in the CERN infrastructure

Coming soon: deployment information

Clusters

666

Nodes

3306

Cores

13370

RAM

31.3 TB

Local Storage

160 TB

Kubernetes

Lingua franca of the cloud

Managed services offered by all major public clouds

Multiple options for on-premise or self-managed deployments

Common declarative API for basic infrastructure : compute, storage, networking

Healthy ecosystem of tools offering extended functionality



Google Cloud Platform



OPENSIFT



Rancher
Kubernetes Engine



MAGNUM



an OpenStack Community Project



docker



containerd



Prometheus



fluentd



envoy

Public Cloud Motivation

Cover for **periodic load spikes**

Access to a (much) larger amount of otherwise scarce resources: **GPUs**

Access resources not available on premises: **TPUs, IPUs**

Evaluate resources in advance, prepare for tenders: GPUs, ARM, ...

Disaster Recovery

Previous Work (2019)

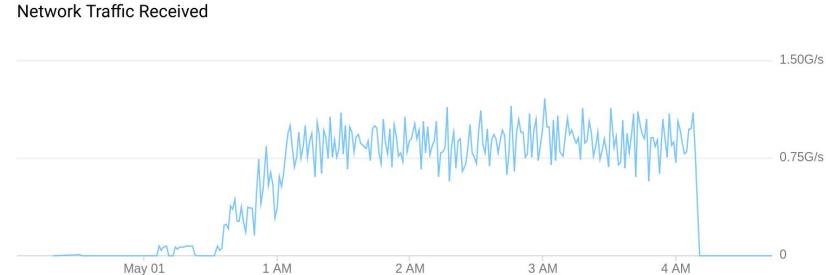
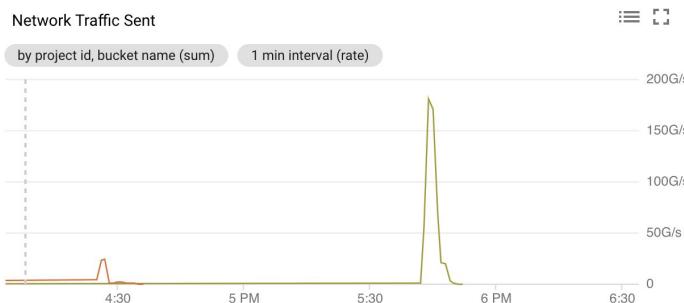
Kubecon Higgs Analysis Demo: **70TB in 5min**

Deep Dive: <https://indico.cern.ch/event/828111/>

Throughput up to **32Gb per VM**

In memory storage used to handle this

Peak at close to **200GB/s**



Google Cloud



70 TB Dataset



Cluster on GKE

Max 25000 Cores

Single Region, 3 Zones

25000 Kubernetes Jobs



Job Results



jupyter
Interactive Visualization

Aggregation

<https://cloud.google.com/customers/cern>

The screenshot shows a video conference interface. On the left, a window displays a histogram of particle energy distribution with a peak around 125 GeV, labeled "CMS Open Data". On the right, a man in a dark hoodie is speaking on stage at a conference. Below the video, the stage has "KubeCon + CloudNativeCon Europe 2019" logos. In the foreground, a web browser window shows the "KubeCon + CloudNativeCon Europe 2019" website. The website features a diagram illustrating the data processing pipeline: "Event Data" leads to "20k+ Core Kits Clusters", which then produce "Summary Data" and a "Real Data" histogram. Text on the site includes "70 TB of Physics Data" and "~25000 Files". The bottom of the browser window says "Live video streaming brought to you by Google Cloud".

<https://www.youtube.com/watch?v=CTfp2woVEkA>

The image shows the CERN Openlab landing page. At the top, a large banner features a photograph of the Large Hadron Collider (LHC) and the text "Helping researchers at CERN to analyze powerful data and uncover the secrets of our universe". Below the banner, a section titled "Google Cloud results" discusses how Google Cloud sped up terabyte-size workloads. A separate section titled "Researchers analyze 70 TB Higgs boson data in minutes" highlights the analysis of 70 TB of Higgs boson data. To the right, a detailed CMS histogram titled "0.0 fb⁻¹ (7 TeV), 0.0 fb⁻¹ (8 TeV)" plots "Events / 3 GeV" against "m_{4l} (GeV)". The legend identifies several data series: "Data" (black dots), "m_H = 125 GeV" (red shaded area), "ZZ → 4l" (blue shaded area), "Zγ* + X" (green shaded area), and "t̄t" (grey shaded area). The x-axis ranges from 80 to 180 GeV, and the y-axis ranges from 0 to 30 events.

But this got us less than halfway...

Support multiple cloud providers, regions

Integrate with our usual in house end user tools

Centralized monitoring, logging, tracing is a requirement

Accounting, costing at a fine grain level is another

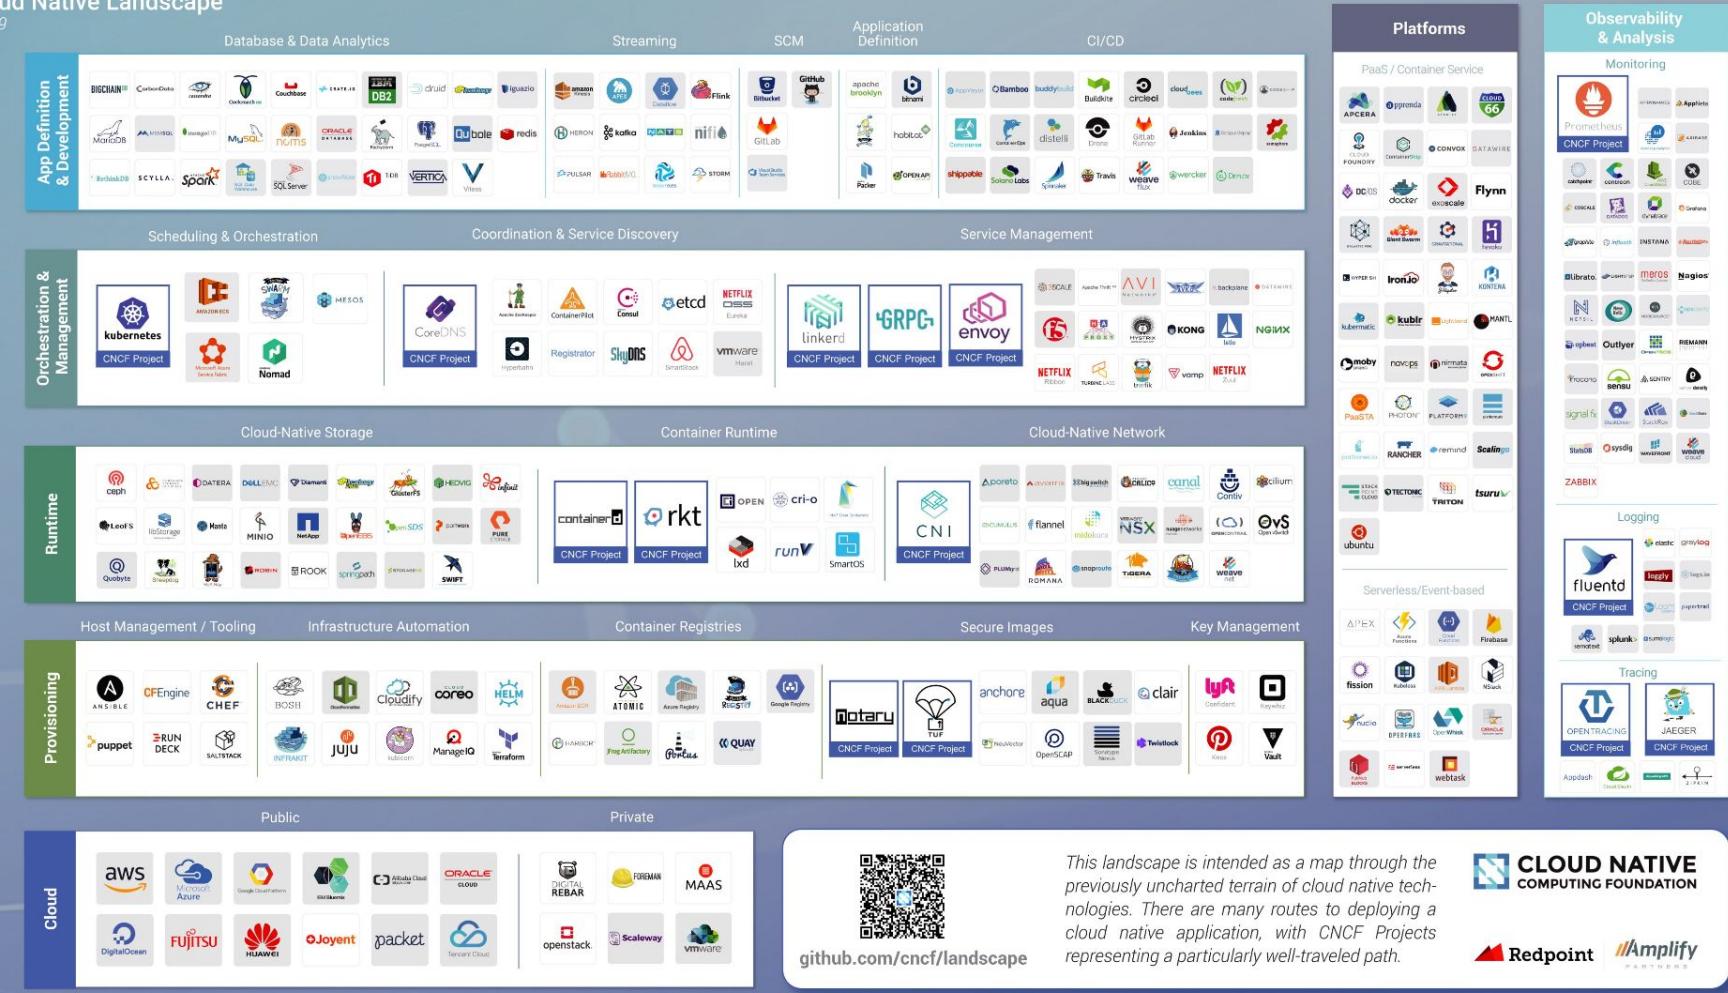
... add here , and here , and here ...

Cloud Native

“” Cloud native technologies empower organizations to build and run scalable applications in modern, dynamic environments such as public, private, and hybrid clouds. Containers, service meshes, microservices, immutable infrastructure, and declarative APIs exemplify this approach.

These techniques enable **loosely coupled systems that are resilient**, manageable, and observable. Combined with robust automation, they allow engineers to make high-impact changes frequently and predictably with minimal toil. “”

Cloud Native Landscape

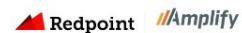


This landscape is intended as a map through the previously uncharted terrain of cloud native technologies. There are many routes to deploying a cloud native application, with CNCF Projects representing a particularly well-traveled path.



github.com/cncf/landscape

Greved logos are not open source



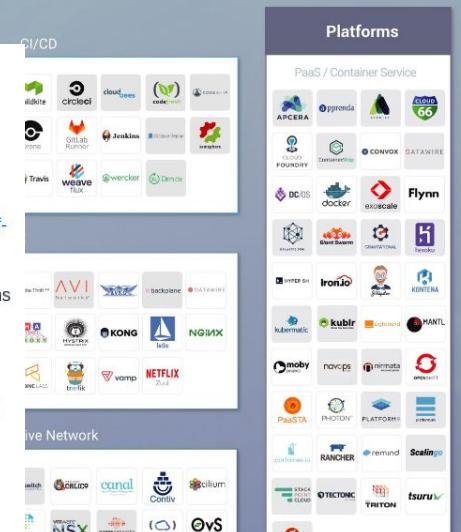


Incubating Stage

Note: The incubation level is the point at which we expect to perform full **due diligence** on projects.

To be accepted to incubating stage, a project must meet the sandbox stage requirements plus:

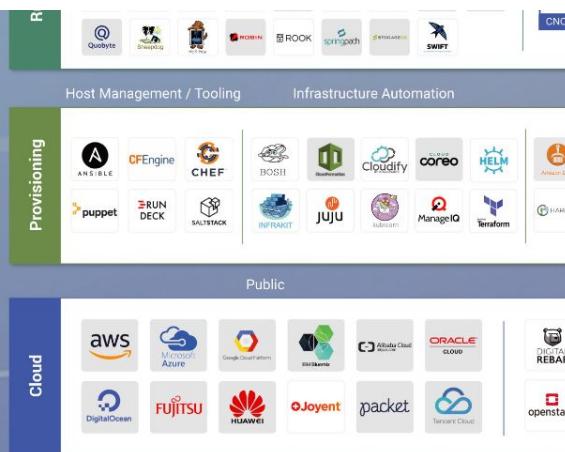
- Document that it is being used successfully in production by at least three independent end users which, in the TOC's judgement, are of adequate quality and scope. For the definition of an end user, see <https://github.com/cncf/toc/blob/main/FAQ.md#what-is-the-definition-of-an-end-user>
- Have a healthy number of committers. A committer is defined as someone with the commit bit; i.e., someone who can accept contributions to some or all of the project.
- Demonstrate a substantial ongoing flow of commits and merged contributions.
- Since these metrics can vary significantly depending on the type, scope and size of a project, the TOC has final judgement over the level of activity that is adequate to meet these criteria
- A clear versioning scheme.
- Specifications must have at least one public reference implementation.



Graduation Stage

To graduate from sandbox or incubating status, or for a new project to join as a graduated project, a project must meet the incubating stage criteria plus:

- Have committers from at least two organizations.
- Have achieved and maintained a Core Infrastructure Initiative [Best Practices Badge](#).
- **Have completed an independent and third party security audit** with results published of similar scope and quality as the following example (including critical vulnerabilities addressed): <https://github.com/envoyproxy/envoy#security-audit> and all critical vulnerabilities need to be addressed before graduation.
- Explicitly define a project governance and committer process. This preferably is laid out in a GOVERNANCE.md file and references an OWNERS.md file showing the current and emeritus committers.
- Have a public list of project adopters for at least the primary repo (e.g., ADOPTERS.md or logos on the project website). For a specification, have a list of adopters for the implementation(s) of the spec.



Deployment

Deployments

Automate and manage centrally full deployment

Clusters, base components, end user applications

Support clusters in **multiple clouds, regions, zones**

Include **base services** for monitoring, *tracing*, policy, software distribution, cost, ...

Support multiple ways to define workloads / services: **helm, kustomize, plain, ...**

Deployment

Includes 3 groups of components

Underlay, handles the cluster(s) deployment and lifecycle

Infrastructure, handles the basic infra services running on all clusters

Services / Applications, the actual workloads running in one or more clusters

All managed in one single git repo and relying on **ArgoCD**

Plain yaml, Helm, Kustomize

Hides a lot of the quite complex deployment details

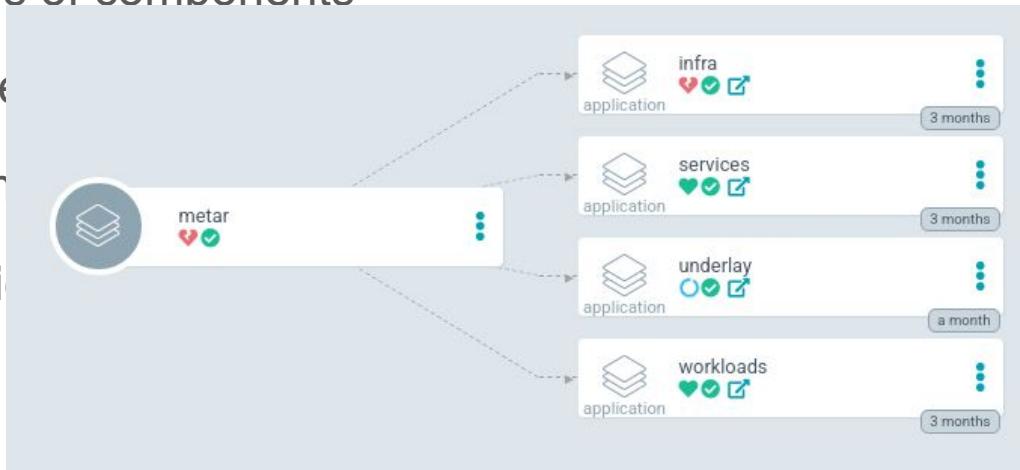
Deployment

Includes 3 groups of components

Underlay, handles

Infrastructure, han

Services / Appli



isters

more clusters

All managed in one single git repo and relying on **ArgoCD**

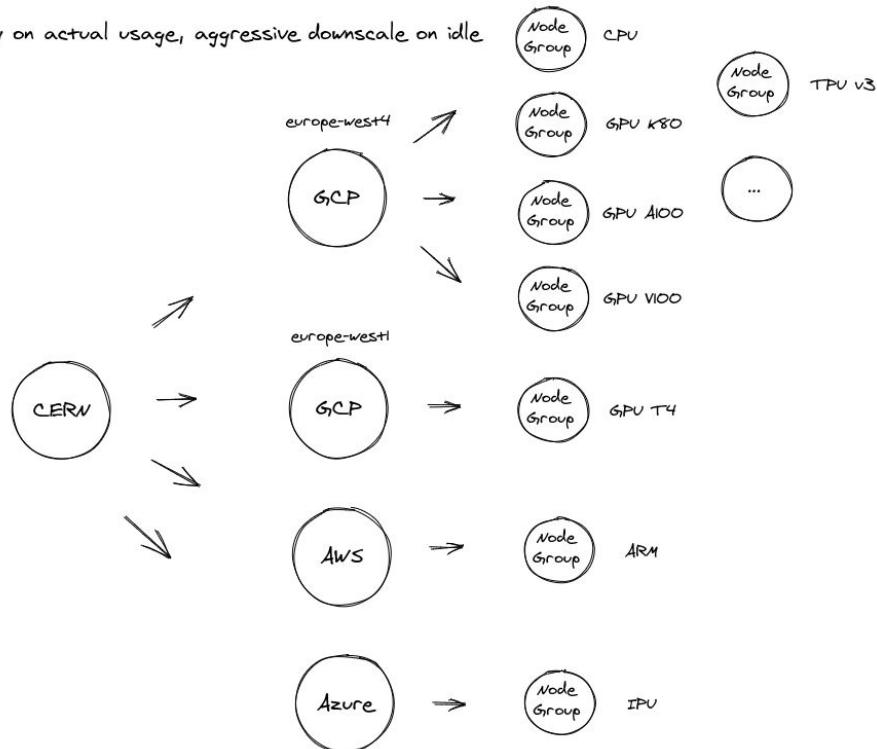
Plain yaml, Helm, Kustomize

Hides a lot of the quite complex deployment details

Deployment: Underlay

All node groups auto scaling on demand

Pay only on actual usage, aggressive downscale on idle



```
clusters:
  gke-europe-west4-a-1:
    cloud: gcp
    autoprovded: true
    config: gcpconfig
    providerConfig:
      location: "europe-west4-a"
      enableTpu: true
      initialClusterVersion: "1.18.12-gke.1210"
    pools:
      default:
        diskSize: 120
        diskType: pd-ssd
        machineType: n1-standard-4
        nodeCount: 2
      a100:
        accelerator:
          count: 1
          type: nvidia-tesla-a100
        autoscaling:
          enabled: true
          minCount: 0
          maxCount: 10
        diskSize: 120
        diskType: pd-ssd
        machineType: a2-highgpu-1g
        nodeCount: 0
        preemptible: true
```

Deployment: Underlay

All node groups auto scaling on demand

Pay only on actual usage, aggressive downscale on idle

```
apiVersion: container.gcp.crossplane.io/v1beta1
kind: GKECluster
metadata:
  name: {{ $cluster }}
spec:
  providerConfigRef:
    name: {{ $clusterdata.config }}
  writeConnectionSecretToRef:
    name: {{ $cluster}}-secret
  namespace: default
```



```
apiVersion: container.gcp.crossplane.io/v1alpha1
kind: NodePool
metadata:
  name: {{ $cluster }}-{{ $pool }}
spec:
  providerConfigRef:
    name: {{ $clusterdata.config }}
  forProvider:
    clusterRef:
      name: {{ $cluster }}
    locations:
      - '{{ $clusterdata.location }}'
  initialNodeCount: {{ $pooldata.nodeCount }}
  {{- if $pooldata.autoscaling }}
  autoscaling:
    autoprovded: {{ $pooldata.autoscaling.autoprovded | default "false" }}
    enabled: {{ $pooldata.autoscaling.enabled | default "false" }}
    maxNodeCount: {{ $pooldata.autoscaling.maxCount }}
    minNodeCount: {{ $pooldata.autoscaling.minCount }}
  {{- end }}
  config:
    {{- if $pooldata.accelerator }}
    accelerators:
      - acceleratorCount: {{ $pooldata.accelerator.count }}
        acceleratorType: {{ $pooldata.accelerator.type }}
    {{- end }}
    diskSizeGb: {{ $pooldata.diskSize }}
    diskType: {{ $pooldata.diskType }}
    imageType: {{ $pooldata.imageType | default "cos_containerd" }}
    {{- if $pooldata.labels }}
    labels:
      {{ $pooldata.labels }}
    {{- end }}
    machineType: {{ $pooldata.machineType }}
```

Compose cloud infrastructure and services into custom platform APIs

Crossplane is an open source Kubernetes add-on that enables platform teams to assemble infrastructure from multiple vendors, and expose higher level self-service APIs for application teams to consume, without having to write any code.

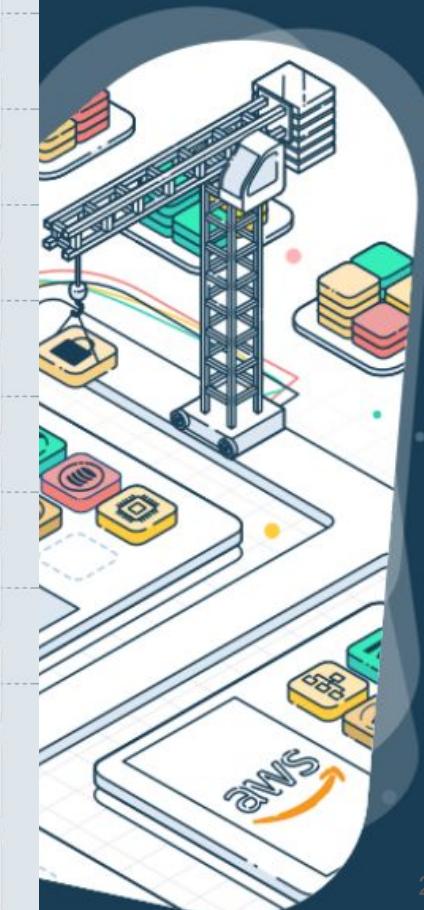
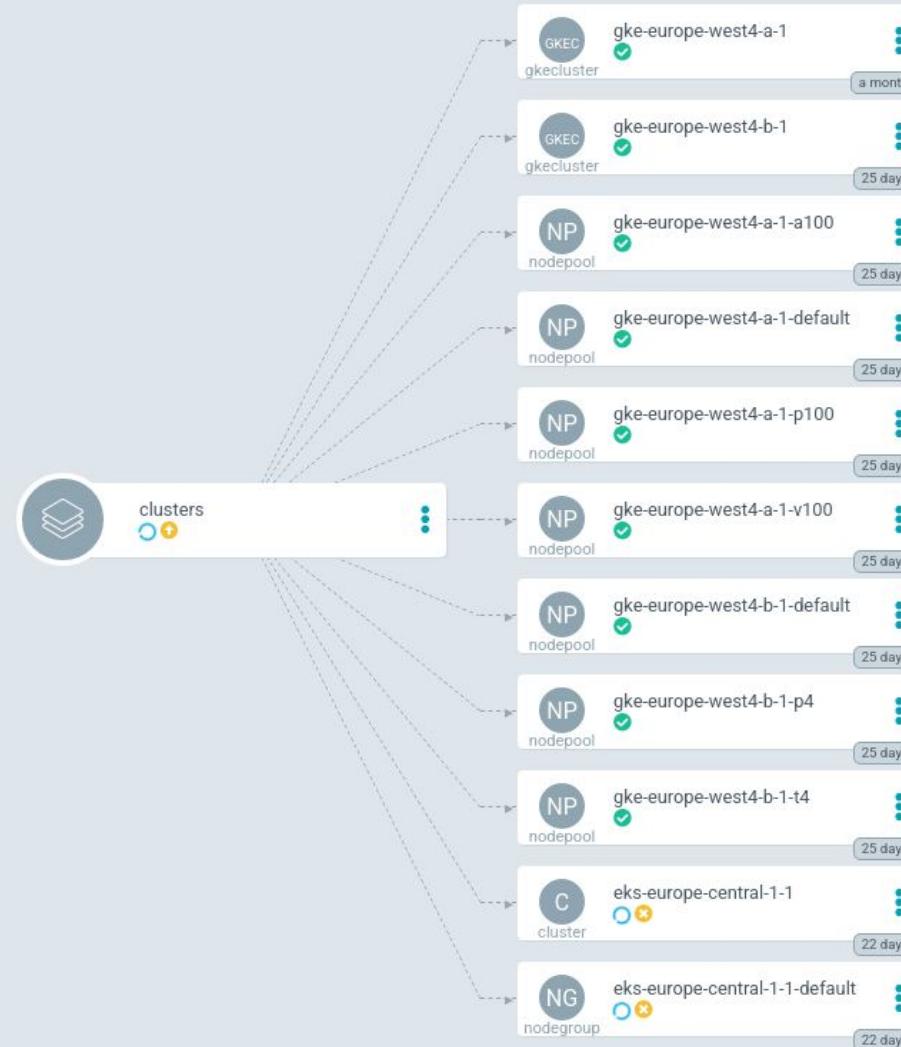
[Get Started on GitHub](#)[Join Slack Channel](#)

Compose cloud infrastructure services into custom platform APIs

Crossplane is an open source tool that enables platform teams to compose cloud infrastructure from multiple providers using higher level self-service APIs, without having to write complex code.

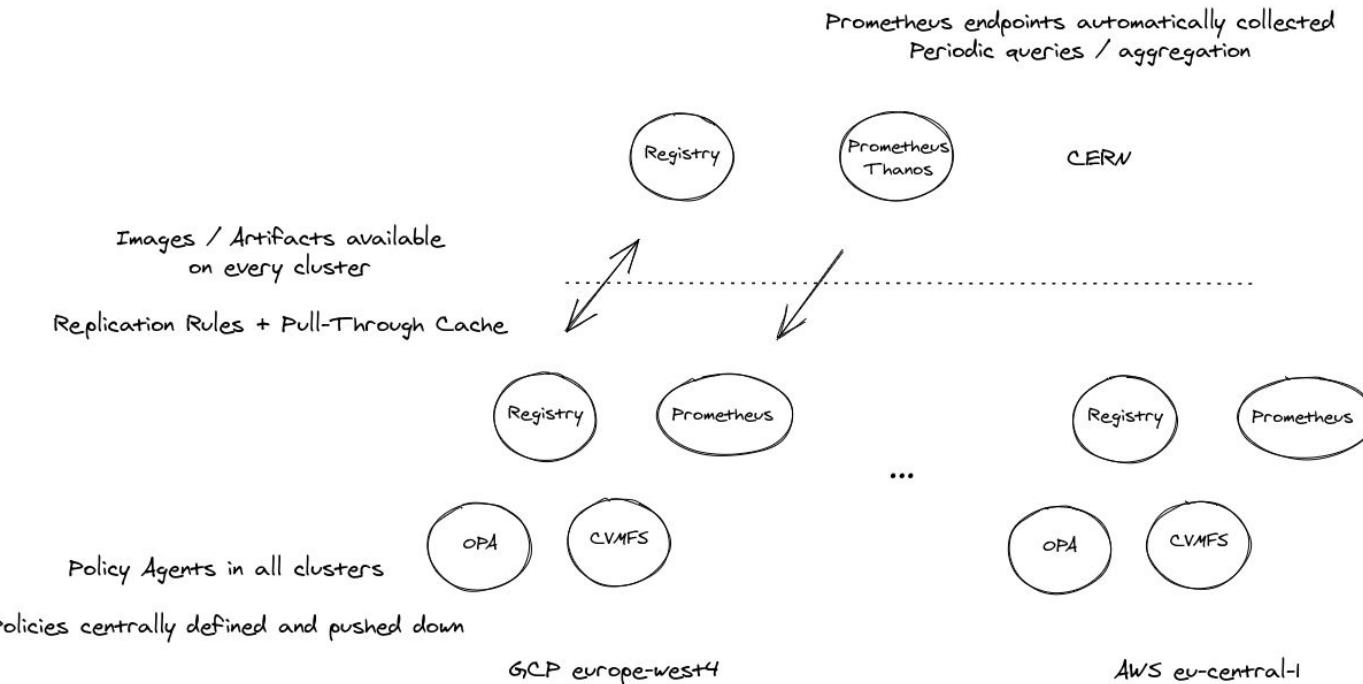


Get Started on GitHub



Deployment: Infrastructure

Base services automatically deployed on all clusters



Use Cases

Use Cases: Accelerated Workloads

Measure performance across a wide range of GPUs

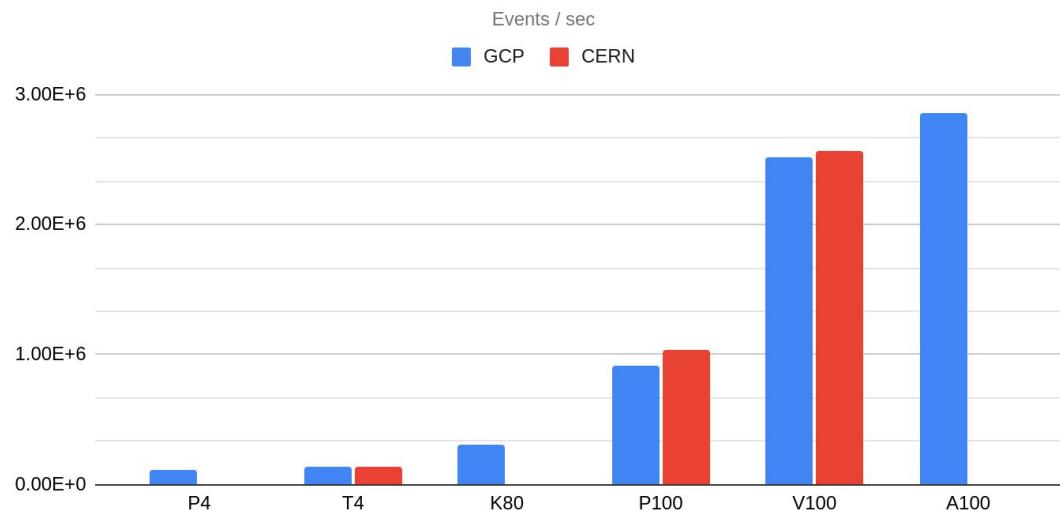
Example with generation of simulation data



Use Cases: Accelerated Workloads

Measure performance across a wide range of GPUs

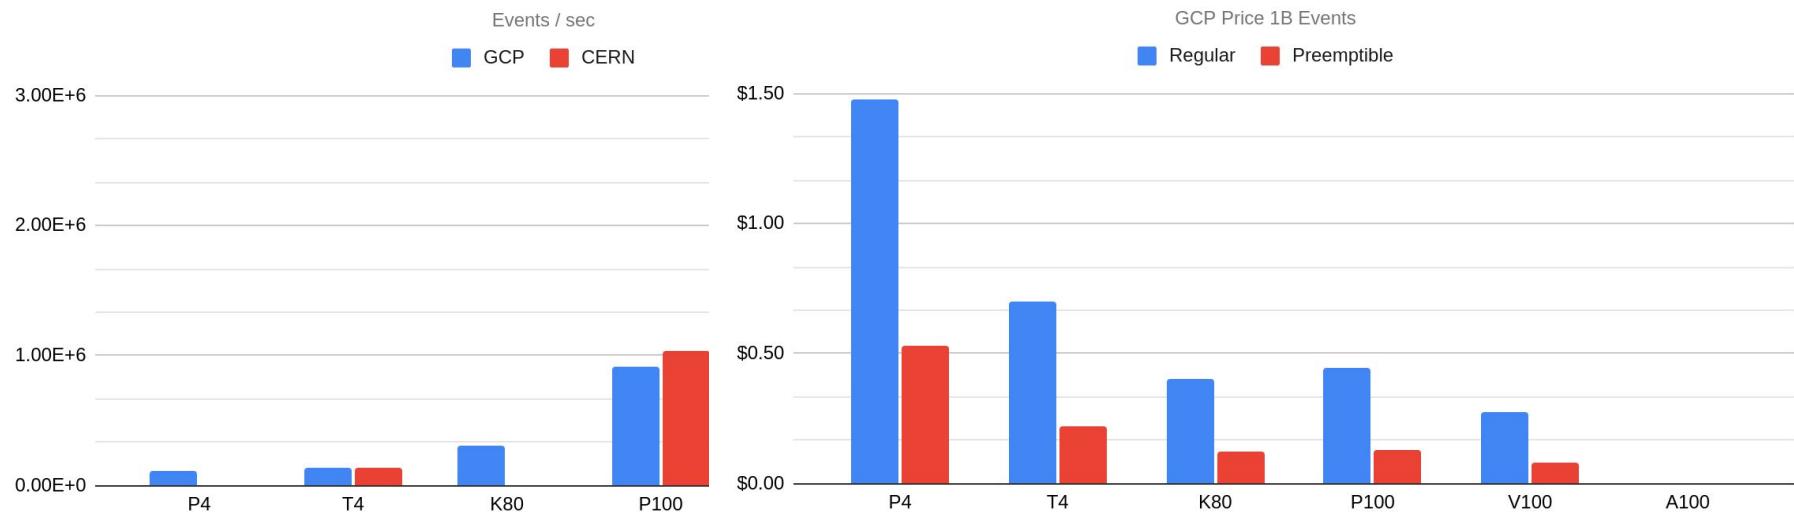
Example with generation of simulation data



Use Cases: Accelerated Workloads

Measure performance across a wide range of GPUs

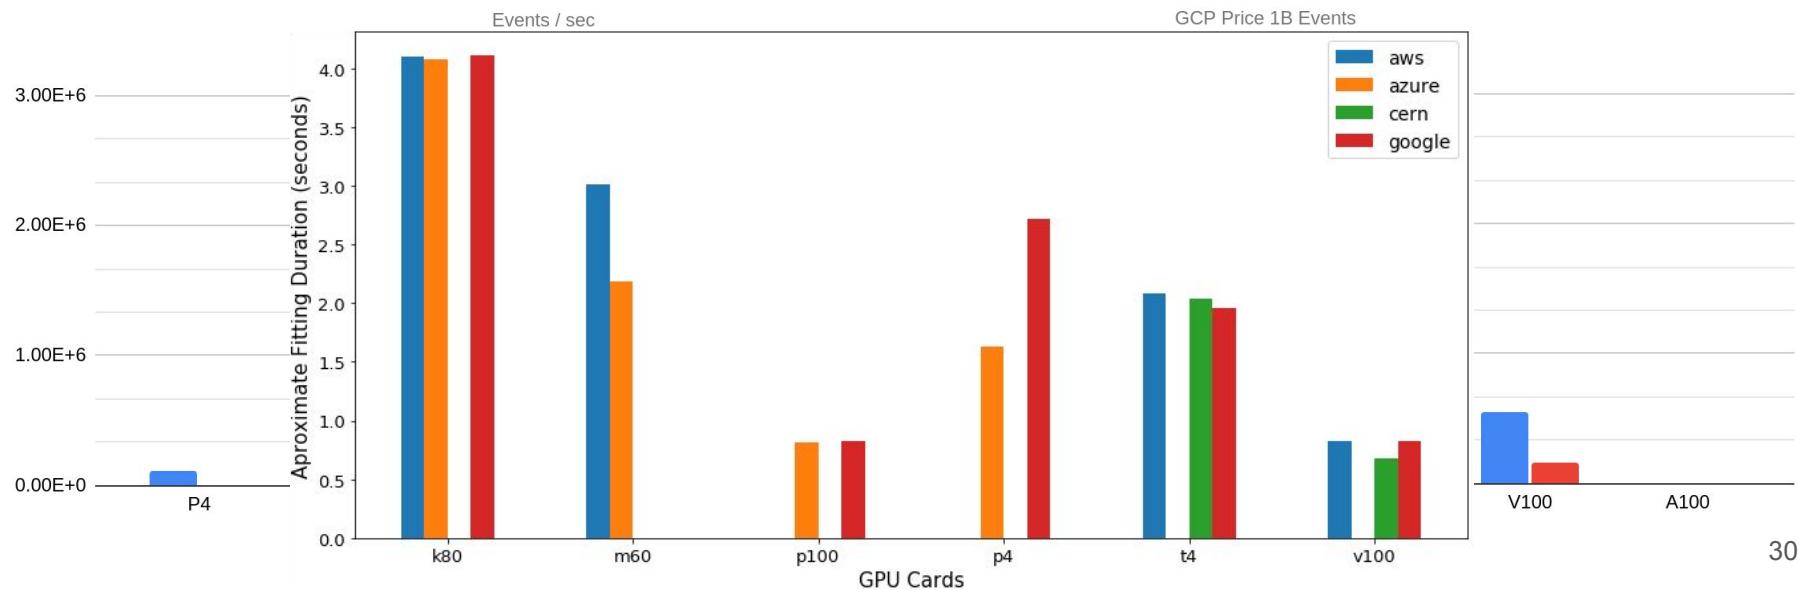
Example with generation of simulation data



Use Cases: Accelerated Workloads

Measure performance across a wide range of GPUs

Example with generation of simulation data



Use Cases: GitLab Runners

<https://gitlab.cern.ch/rbritoda/gitlab-runner-public>

Experimental setup with runners using public cloud resources

Cluster(s) auto scaling on demand

Slight delay (<2min) for job start when a fresh node is required

Large numbers and types of resources available

GPUs: Nvidia A100, V100, P100, K80, P4, T4, M60

TPUs: v2-8, v3-8, can add larger instances if needed

ARM: Graviton2 (m6g.xlarge, m6g.16xlarge, m6g.metal)

| Status | Job ID | Name | Coverage |
|-------------|---|------------------------|-----------------------------|
| Test | | | |
| | #12061346 docker-arm | sample-arm | ⌚ 00:00:05 ⌚ 2 hours ago |
| | #12061347 docker-arm-graviton2 | sample-arm-graviton2 | ⌚ 00:00:08 ⌚ 2 hours ago |
| | #12061338 docker-gpu | sample-gpu | ⌚ 00:00:05 ⌚ 2 hours ago |
| | #12061339 docker-gpu-nvidia | sample-gpu-nvidia | ⌚ 00:00:15 ⌚ 2 hours ago |
| | #12061340 docker-gpu-a100 | sample-gpu-nvidia-a100 | ⌚ 00:00:15 ⌚ 2 hours ago |
| | #12061345 docker-gpu-k80 | sample-gpu-nvidia-k80 | ⌚ 00:00:06 ⌚ 2 hours ago |
| | #12061342 docker-gpu-p100 | sample-gpu-nvidia-p100 | ⌚ 00:00:05 ⌚ 2 hours ago |
| | #12061344 docker-gpu-p4 | sample-gpu-nvidia-p4 | ⌚ 00:00:12 ⌚ 2 hours ago |
| | #12061343 docker-gpu-t4 | sample-gpu-nvidia-t4 | ⌚ 00:00:06 ⌚ 2 hours ago |
| | #12061341 docker-gpu-v100 | sample-gpu-nvidia-v100 | ⌚ 00:00:12 ⌚ 2 hours ago |
| | #12061348 docker-tpu | sample-tpu | ⌚ 00:03:03 ⌚ 2 hours ago |
| | #12061349 docker-tpu-v2-8 | sample-tpu-v2-8 | ⌚ 00:03:01 ⌚ 2 hours ago |
| | #12061350 docker-tpu-v3-8 | sample-tpu-v3-8 | ⌚ 00:02:58 ⌚ 2 hours ago |

Pipeline Needs **Jobs 13** Tests 0

```
62 Waiting for pod gitlab-runner/runner-hdqec7c-project-108952-concurrent-1fbmnv to be running, status is Pending
63     ContainersNotReady: "containers with unready status: [build helper]"
64     ContainersNotReady: "containers with unready status: [build helper]"
65 Running on runner-hdqec7c-project-108952-concurrent-1fbmnv via gitlab-runner-tpu-v3-8-gitlab-runner-5d54df77b-wcqg8...
66 Getting source from Git repository
67 Fetching changes with git depth set to 50...
68 Initialized empty Git repository in /builds/rbritoda/gitlab-runner-public/.git/
69 Created fresh repository.
70 Checking out 828259f8 as load...
71 Skipping Git submodules setup
72 Executing "step script" stage of the job script
73
74 #!/usr/bin/python3
75 import tensorflow as tf
76 import os
77 from tensorflow.python.profiler import profiler_client
78
79 endpoint = os.environ['KUBE_GOOGLE_CLOUD_TPU_ENDPOINTS']
80 print("Connecting to TPU at %s\n" % endpoint)
81
82 tpu = tf.distribute.cluster_resolver.TPUClusterResolver()
83 print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
84
85 tf.config.experimental_connect_to_cluster(tpu)
86 tf.tpu.experimental.initialize_tpu_system(tpu)
87
88 print(profiler_client.monitor(endpoint.replace('8470', '8466'), 100, 2))
89
90 2021-02-09 21:01:12.702455: I tensorflow/core/distributed_runtime/rpc/grpc_server_lib.cc:390] Started server with target: grpc://localhost:30018
91 Connecting to TPU at grpc://10.116.18.178:8470
92 Running on TPU  ['10.116.18.178:8470']
93
94 Timestamp: 21:01:22
95 TPU type: TPU v3
96 Utilization of TPU Matrix Units (higher is better): 0.000%
97 Cleaning up file based variables
98 Job succeeded
```

00:01 00:12 00:00 00:00

```
13 Running on runner-yawjis17o-project-108952-concurrent-18fffd via gitlab-runner-arm-graviton2-gitlab-runner-575dd7ff87-tlpbc...
```

```
15 Getting source from Git repository
```

```
16 Fetching changes with git depth set to 50...
```

```
17 Initialized empty Git repository in /builds/rbritoda/gitlab-runner-public/.git/
```

```
18 Created fresh repository.
```

```
19 Checking out 828259f8 as load...
```

```
20 Skipping Git submodules setup
```

```
22 Executing "step_script" stage of the job script
```

```
23 $ lscpu
```

```
24 Architecture: aarch64
```

```
25 CPU op-mode(s): 32-bit, 64-bit
```

```
26 Byte Order: Little Endian
```

```
27 CPU(s): 2
```

```
28 On-line CPU(s) list: 0,1
```

```
29 Thread(s) per core: 1
```

```
30 Core(s) per socket: 2
```

```
31 Socket(s): 1
```

```
32 NUMA node(s): 1
```

```
33 Vendor ID: ARM
```

```
34 Model: 1
```

```
35 Model name: Neoverse-N1
```

```
36 Stepping: r3p1
```

```
37 BogoMIPS: 243.75
```

```
38 L1d cache: 128 KiB
```

```
39 L1i cache: 128 KiB
```

```
40 L2 cache: 2 MiB
```

```
41 L3 cache: 32 MiB
```

```
42 NUMA node0 CPU(s): 0,1
```

```
43 Vulnerability Itlb multihit: Not affected
```

```
44 Vulnerability Llft: Not affected
```

```
45 Vulnerability Mds: Not affected
```

```
46 Vulnerability Meltdown: Not affected
```

```
47 Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via pr
```

```
48 Vulnerability Spectre v1: Mitigation; __user pointer sanitization
```

```
49 Vulnerability Spectre v2: Not affected
```

```
50 Vulnerability Srbds: Not affected
```

```
51 Vulnerability Tsx async abort: Not affected
```

```
52 Flags:
```

```
fp asimd evtstrm aes pmull sha1 sha2 crc32 atomics tppp asimdrdp cpuid asimdrdm lrcpc dcpop asimddp ssbs
```

```
54 Cleaning up file based variables
```

```
56 Job succeeded
```

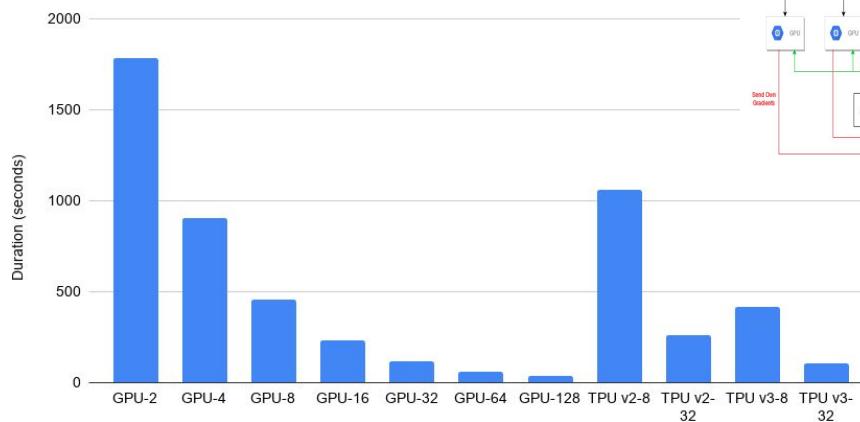


Use Cases: Machine Learning / Kubeflow

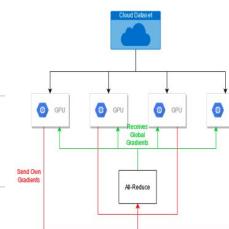
Scale out / distributed training

Example: Fast Simulation with 3D GANs (TensorFlow based)

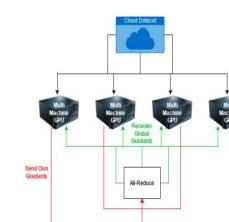
First real attempt at scaling TPUs



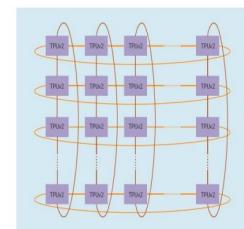
Mirrored Strategy



Multi Worker Mirrored Strategy

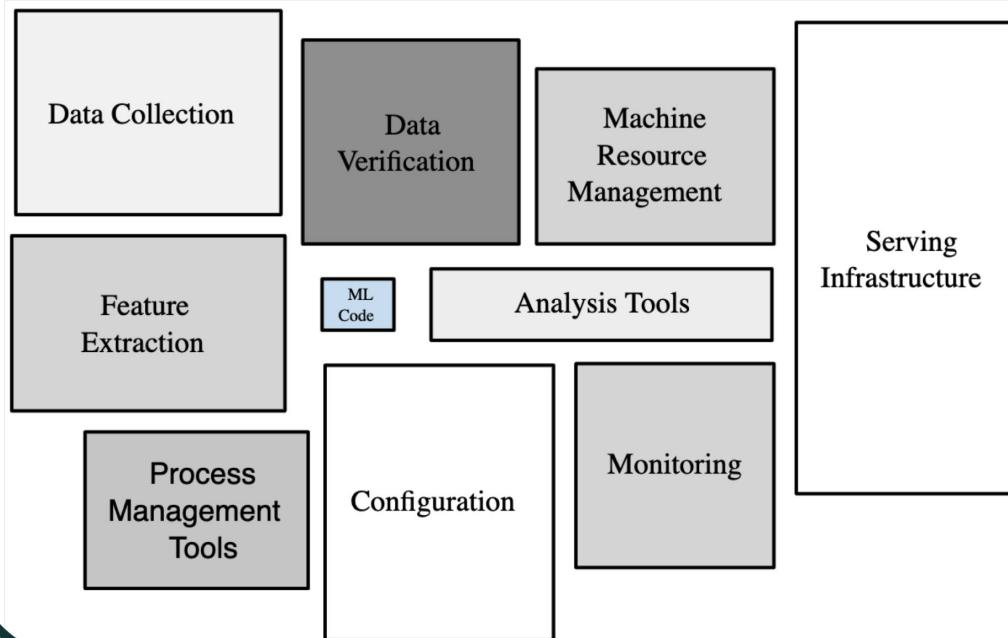


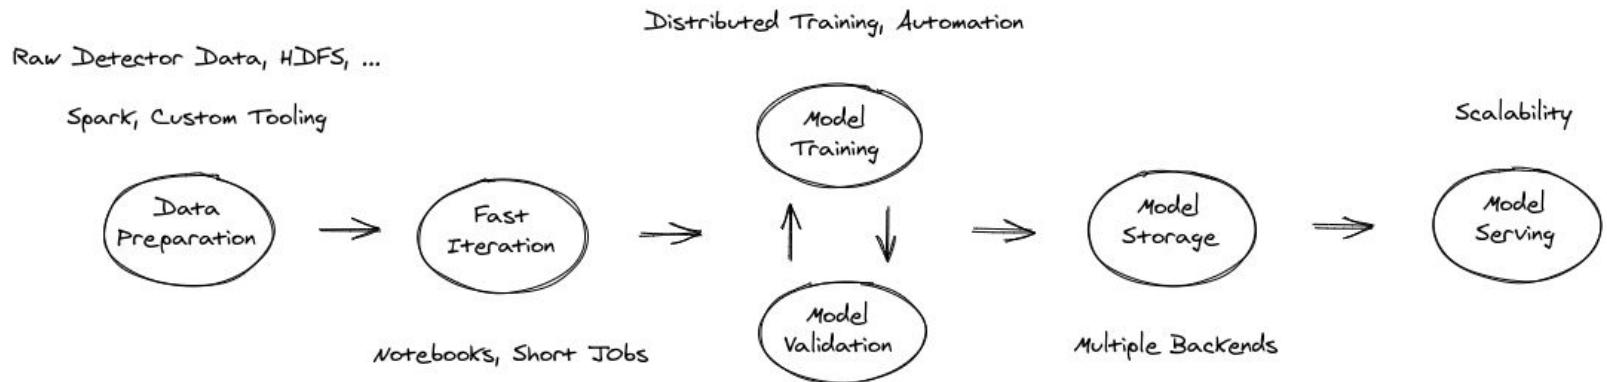
TPU Strategy

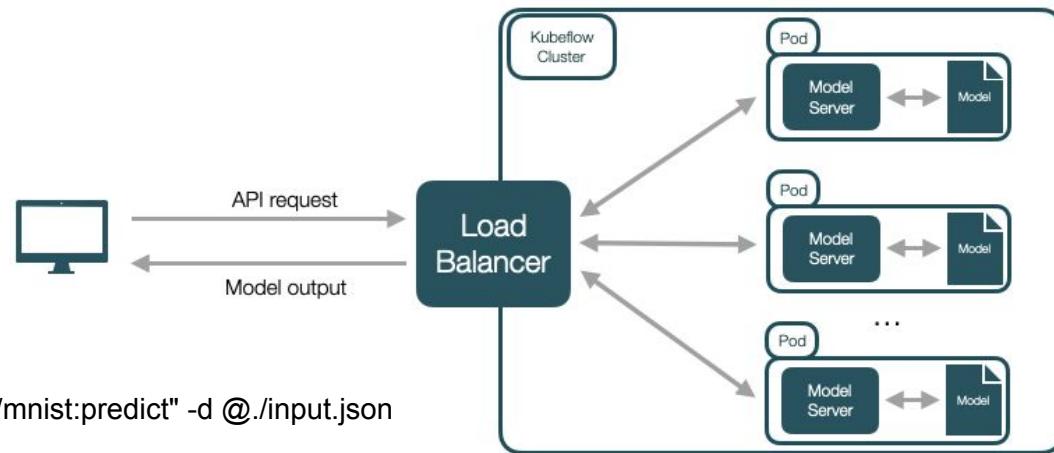
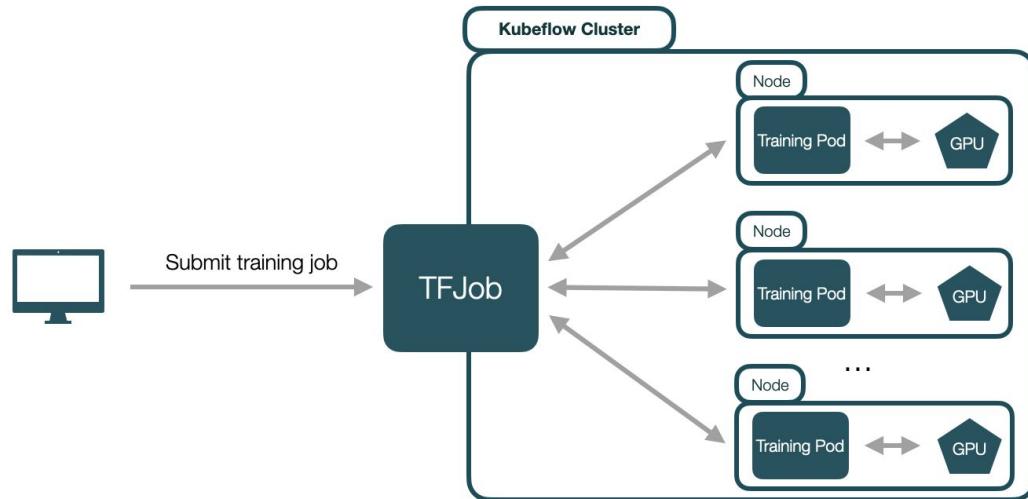




Kubeflow







Inference: curl -v -H "Host: host" "http://host_ip/v1/models/mnist:predict" -d @./input.json

On resources and cost...

Resources are not infinite, even in the public cloud

CPUs kind of feel like infinite, at least at O(10k) scale

GPUs (preemptibles) need more care at O(1k)

Pre-emptibles / spot can be very cost effective

Up to 90% cheaper, a lot (most?) of the workloads we have shine on these

All possible, but **need to stay flexible** on clouds, regions, zones of choice

On resources and cost...

Beware of straight, definite answers

“It’s X times more expensive than on-premises”

X can be anything from 3, to the most common 6, to 20, or 600

Better Answer: It's complicated

Workload dependent, negotiable, small changes give big gains, ...

Pre-emptibles / spot, Sustained / committed use discounts, egress deals, ...

On resources and cost...

Taking our ML example presented earlier...

Running one epoch on 1 GPU takes ~3550 seconds (~1h)

Reserved Cost: \$2.53 USD , Preemptible cost : \$0.73 USD

100 epochs would take over 4 days: 355000 seconds

Workload particularly well behaved, scales almost linearly without much effort

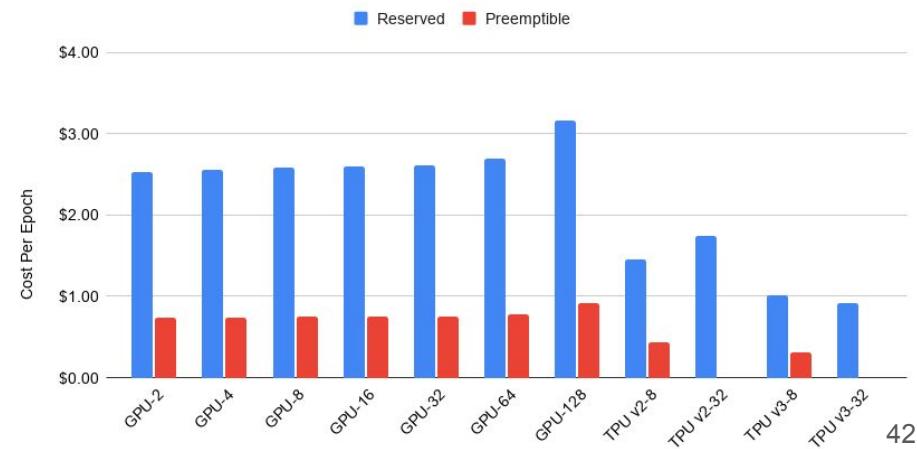
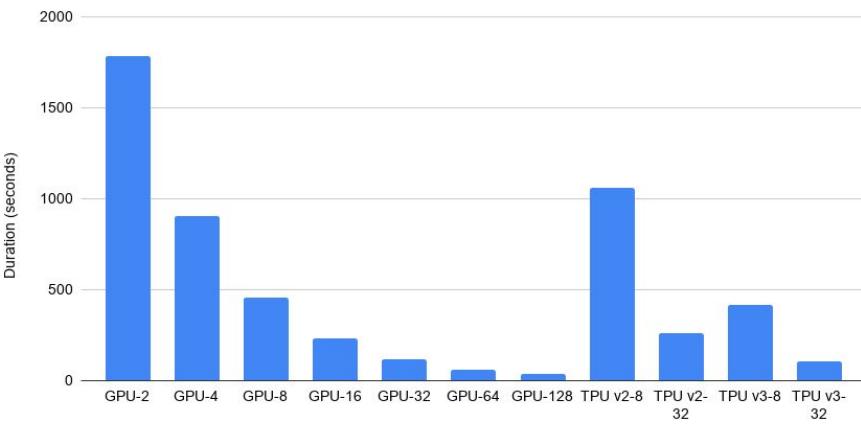
Larger models ~~might not~~ won't scale the same way

On resources and cost...

From 1 to 128 GPUs: 3550 to 35 seconds per epoch

x100 speedup at the same total cost

TPUs seem to be particularly cost effective

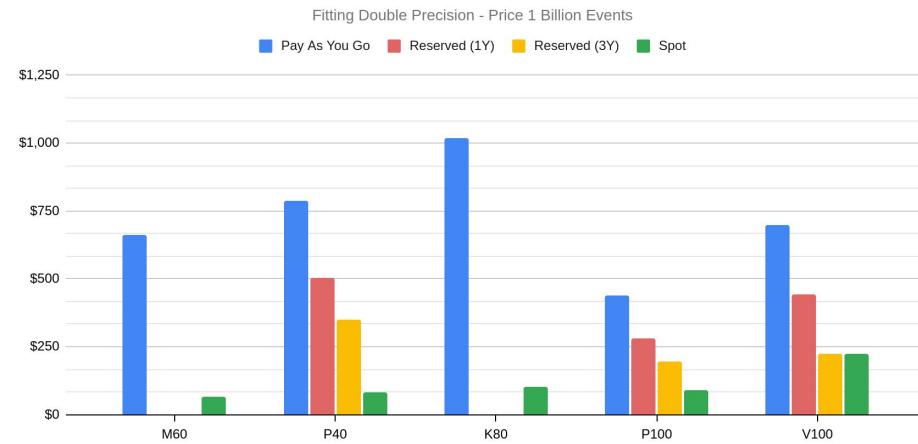
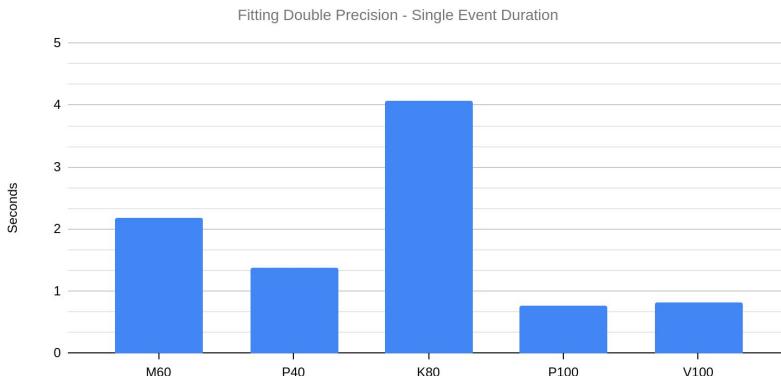


On resources and cost...

Taking another example doing statistical inference

The most powerful GPU does not give the best *value for money* (far from it)

Workloads fitting the spot model shine in the public cloud



Billing down to individual workloads(s) →

The screenshot shows the kubecost UI for a GKE cluster named 'gke-kubecost-demo-001'. The main dashboard displays monthly savings of \$113.47 and monthly cluster costs of \$293.16 with a 41.1% cost efficiency. Below this, a chart shows resource efficiency for Compute, Memory, Storage, and GPU across four categories: Idle, System, Apps, and Other. A pie chart illustrates controller allocation across various namespaces. On the left, a sidebar provides navigation links for Overview, Cost Allocation, Assets, Savings, Health, Reports, and Notifications. The bottom section shows a detailed breakdown of costs by pod over the last 7 days, with a table listing individual workloads and their resource usage and costs.

← **Cost efficiency evaluation**

This screenshot shows a detailed view of cost efficiency evaluation. It includes a bar chart titled 'Last 7 days by pod daily' showing resource usage for pods from April 18 to April 21. A modal window on the right provides a breakdown of costs by pod, including metrics like CPU, GPU, RAM, PV, Network, LB, and Shared. The modal also includes filters for Date Range (Last 7 days), Namespace, Controller kind, Service, Pod, Department, Environment, Owner, Product, and Team. The total cost for the selected period is \$4.51.

| Name | CPU | GPU | RAM | PV | Network | LB | Shared | Efficiency | Total cost |
|---|---------|---------|---------|--------|---------|--------|--------|------------|------------|
| Totals | \$28.29 | \$12.71 | \$14.12 | \$0.63 | \$0.00 | \$8.23 | \$0.00 | 47.3% | \$63.97 |
| __idle__ | \$12.51 | -\$3.14 | \$12.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | — | \$21.37 |
| kubecost-cost-analyzer-7874d5669f-w4sz4 | \$1.24 | \$0.00 | \$0.09 | \$0.32 | \$0.00 | \$4.11 | \$0.00 | \$0.00 | \$5.76 |
| tf-dtraining-78d7456717-2svqp | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.38 |
| tf-dtraining-78d7456717-4n857 | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 |
| tf-dtraining-78d7456717-dc71f | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 |
| cpu-runner-7cd55667bc-d69cv | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 |
| cpu-runner-7cd55667bc-t8wd6 | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 |

Conclusion

Managing **100s of clusters** on-premises, **10s** in multiple public clouds / regions

Cluster launch to full production availability in **less than 15 minutes**

Flexibility helps get to the most cost effective deployment

Kubernetes API is the **common language and enabler**

Questions?



KubeCon



CloudNativeCon

Europe 2022

16 – 20 MAY

VALENCIA, SPAIN + VIRTUAL

REGISTER

EXPLORE THE SCHEDULE

<https://events.linuxfoundation.org/kubecon-cloudnativecon-europe/>

Cost Allocation

Assets

Savings

Health

Reports

Notifications

Overview / gke-kubecost-demo-001



Monthly savings of \$113.47 identified

LEARN MORE >

Monthly cost

\$293.16

Cost Efficiency

41.1%

Monthly cluster costs

Monthly run rate expenses based on resource prices

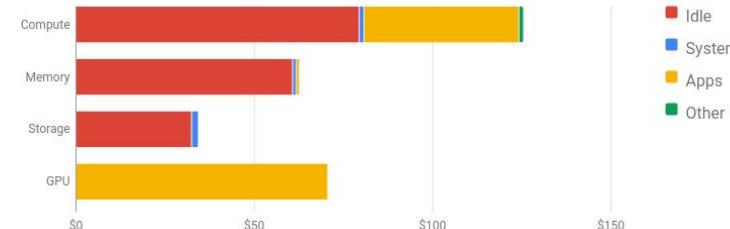


Total cost ▾

CLUSTER METRICS >

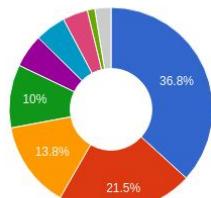
Resource Efficiency

Based on currently provisioned resources and last 7d usage

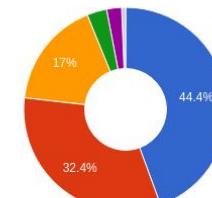


VIEW ASSETS >

Controller Allocation


 kubeflow/deployment:tf-dtr...
 gitlab-runner/deployment:c...
 kube-cost/deployment:kube...
◀ 1/3 ▶

Service Allocation


 kube-cost-cost-analyzer
 kube-grafana
 kube-dns
 kube-prometheus-server
◀ 1/2 ▶
 Switch cluster
gke-kubecost-demo-001

COST ALLOCATION >

Settings

Namespace

Monthly Cost

Efficiency

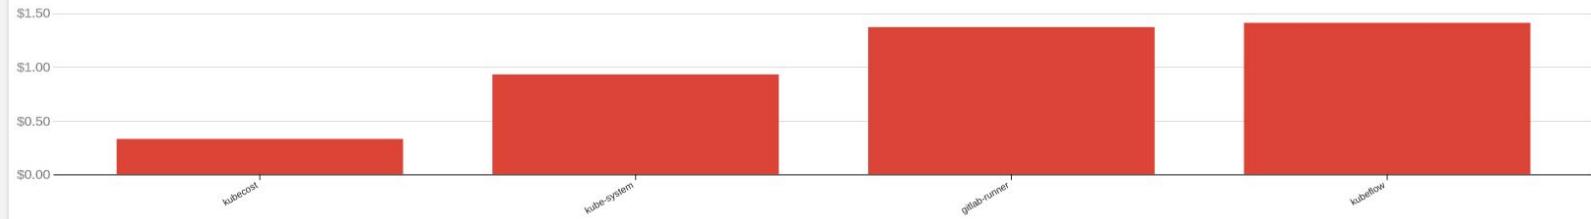
Infrastructure health

 Cost Allocation Assets Savings Health Reports Notifications

 Note: this page is actively being replaced. View [Reports](#) for the new UI.

Total measured cost

Top 25 results measured over time window



| Namespace | Mem Cost | CPU Cost | Efficiency | Network Cost | PV Cost | GPU Cost | External Cost | Total Cost |
|-------------------------------|---------------|---------------|------------|-----------------------------|---------------|---------------|-------------------------|---------------|
| kubeflow | \$0.00 | \$0.03 | 100% | Add network | \$0.00 | \$1.37 | Add Key | \$1.41 |
| gitlab-runner | \$0.15 | \$1.22 | 88.9% | Add network | \$0.00 | \$0.00 | Add Key | \$1.37 |
| kube-system | \$0.11 | \$0.81 | 9.3% | Add network | \$0.00 | \$0.00 | Add Key | \$0.93 |
| kube-cost | \$0.03 | \$0.20 | 25.1% | Add network | \$0.10 | \$0.00 | Add Key | \$0.33 |
| Totals | \$0.29 | \$2.26 | | \$0.00 | \$0.10 | \$1.37 | \$0.00 | \$4.04 |

Assets

18 April 2021 through now by Category

Date Range

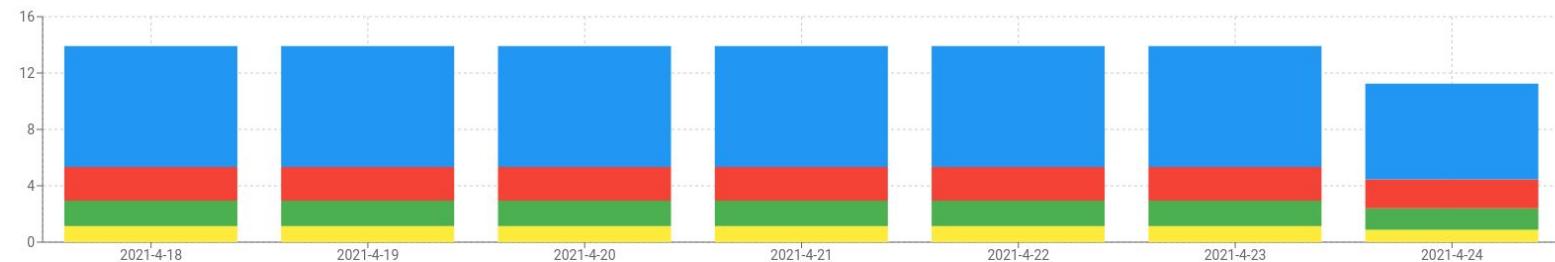
Last 7 days

Resolution

Daily

Breakdown

Category



| Name | Credits | Adjusted | Total cost |
|------------|---------|----------|------------|
| Totals | \$0.00 | \$0.00 | \$94.75 |
| Compute | \$0.00 | \$0.00 | \$58.25 |
| Management | \$0.00 | \$0.00 | \$16.45 |
| Network | \$0.00 | \$0.00 | \$12.33 |
| Storage | \$0.00 | \$0.00 | \$7.72 |

Switch cluster

Rows per page: 25 ▾ 1-4 of 4 < >

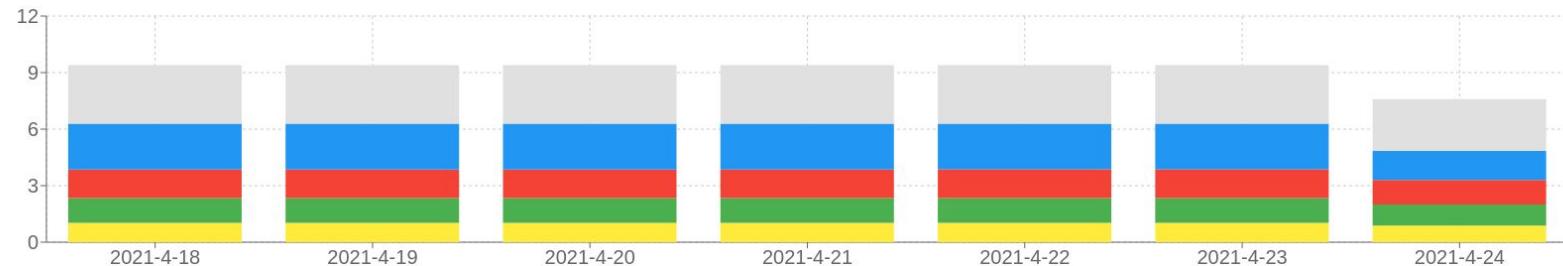
[Overview](#)[Cost Allocation](#)[Assets](#)[Savings](#)[Health](#)[Reports](#)[Notifications](#)

Reports / Last 7 days by namespace daily



Last 7 days by namespace daily

18 April 2021 through now by Namespace

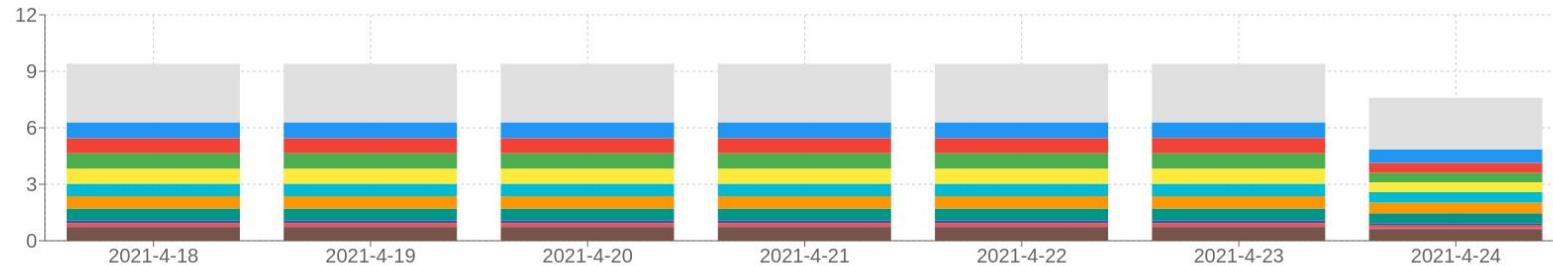


| Name | CPU | GPU | RAM | PV | Network | LB | Shared | External | Efficiency | Total cost |
|---------------------------------|---------|---------|---------|--------|---------|--------|--------|----------|------------|------------|
| Totals | \$28.29 | \$12.71 | \$14.12 | \$0.63 | \$0.00 | \$8.23 | \$0.00 | \$0.00 | 47.3% | \$63.97 |
| ⋮ __idle__ | \$12.51 | -\$3.14 | \$12.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | — | \$21.37 |
| ⋮ kubeflow | \$0.27 | \$15.84 | \$0.01 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$16.12 |
| ⋮ kubecost | \$1.32 | \$0.00 | \$0.21 | \$0.63 | \$0.00 | \$8.23 | \$0.00 | \$0.00 | 34.8% | \$10.40 |
| ⋮ gitlab-runner | \$8.02 | \$0.00 | \$1.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$9.02 |
| ⋮ kube-system | \$6.16 | \$0.00 | \$0.90 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 7.8% | \$7.06 |

[Switch cluster
gke-kubecost-demo-001](#)[Settings](#)

Last 7 days by pod daily

18 April 2021 through now by Pod



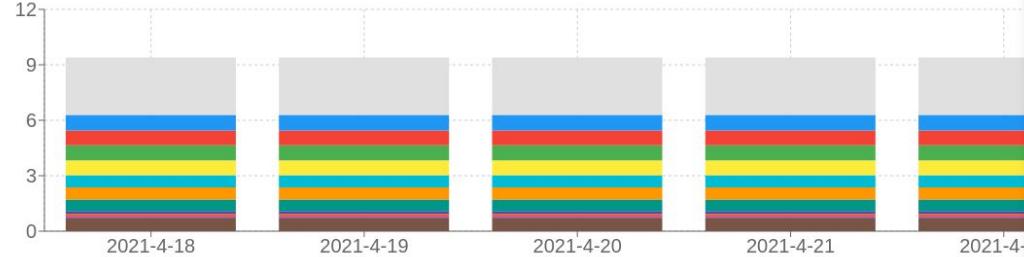
| Name | CPU | GPU | RAM | PV | Network | LB | Shared | External | Efficiency | ↓ | Total cost |
|--|---------|---------|---------|--------|---------|--------|--------|----------|------------|---------|------------|
| Totals | \$28.29 | \$12.71 | \$14.12 | \$0.63 | \$0.00 | \$8.23 | \$0.00 | \$0.00 | 47.3% | \$63.97 | |
| ... __idle__ | \$12.51 | -\$3.14 | \$12.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | — | \$21.37 | |
| ... kube-cost-analyzer-7874d5669fw4sz4 | \$1.24 | \$0.00 | \$0.09 | \$0.32 | \$0.00 | \$4.11 | \$0.00 | \$0.00 | 4.1% | \$5.76 | |
| ... tf-dtraining-78d74567f7-2svqp | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.38 | |
| ... tf-dtraining-78d74567f7-4n857 | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 | |
| ... tf-dtraining-78d74567f7-dc7lf | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 | |
| ... cpu-runner-7cd55667bc-d69cv | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 | |
| ... cpu-runner-7cd55667bc-t8wd6 | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 | |

Switch cluster
gke-kubecost-demo-001

Settings

Last 7 days by pod daily

18 April 2021 through now by Pod



Date Range
Last 7 days

Filters

namespace

e.g. cluster "dev", r

Shared resources

Add a shared name

Share weighted

Cluster

Node

Namespace

Controller kind

Controller

Service

Pod

Department

Environment

Owner

Product

Team

costs

Resolution

Separate

Daily

+

label "app:cost-analyzer"

ost in settings

SETTINGS

| Name | CPU | GPU | RAM | PV | Network | LB | Shared | Efficiency | Total cost |
|--|---------|---------|---------|--------|---------|--------|--------|------------|-------------|
| Totals | \$28.29 | \$12.71 | \$14.12 | \$0.63 | \$0.00 | \$8.23 | \$0.00 | 47.3% | \$63.97 |
| __idle__ | \$12.51 | -\$3.14 | \$12.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | — | \$21.37 |
| kubecost-cost-analyzer-7874d5669fw4sz4 | \$1.24 | \$0.00 | \$0.09 | \$0.32 | \$0.00 | \$4.11 | \$0.00 | \$0.00 | 4.1% \$5.76 |
| tf-dtraining-78d74567f7-2svqp | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.38 |
| tf-dtraining-78d74567f7-4n857 | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 |
| tf-dtraining-78d74567f7-dc7lf | \$0.09 | \$5.28 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | Inf% | \$5.37 |
| cpu-runner-7cd55667bc-d69cv | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 |
| cpu-runner-7cd55667bc-t8wd6 | \$4.01 | \$0.00 | \$0.50 | \$0.00 | \$0.00 | \$0.00 | \$0.00 | 88.9% | \$4.51 |

Switch cluster
gke-kubecost-demo-001

Settings