

Jaime S. Cardoso

Associate Professor

jaime.cardoso@inesctec.pt

<http://www.inescporto.pt/~jsc/>

<http://medicalresearch.inescporto.pt/breastresearch/>

INESC TEC and Faculdade de Engenharia, Universidade do Porto,
Portugal


Deep Learning: Variations on the Theme and Current Trends

Data Science Portugal

Feb 21st, 2017, Portugal



Data Driven Design

- When to use?
 - Difficult to reason about a generic rule that solves the problem
 - Easy to collect examples (with the solution)



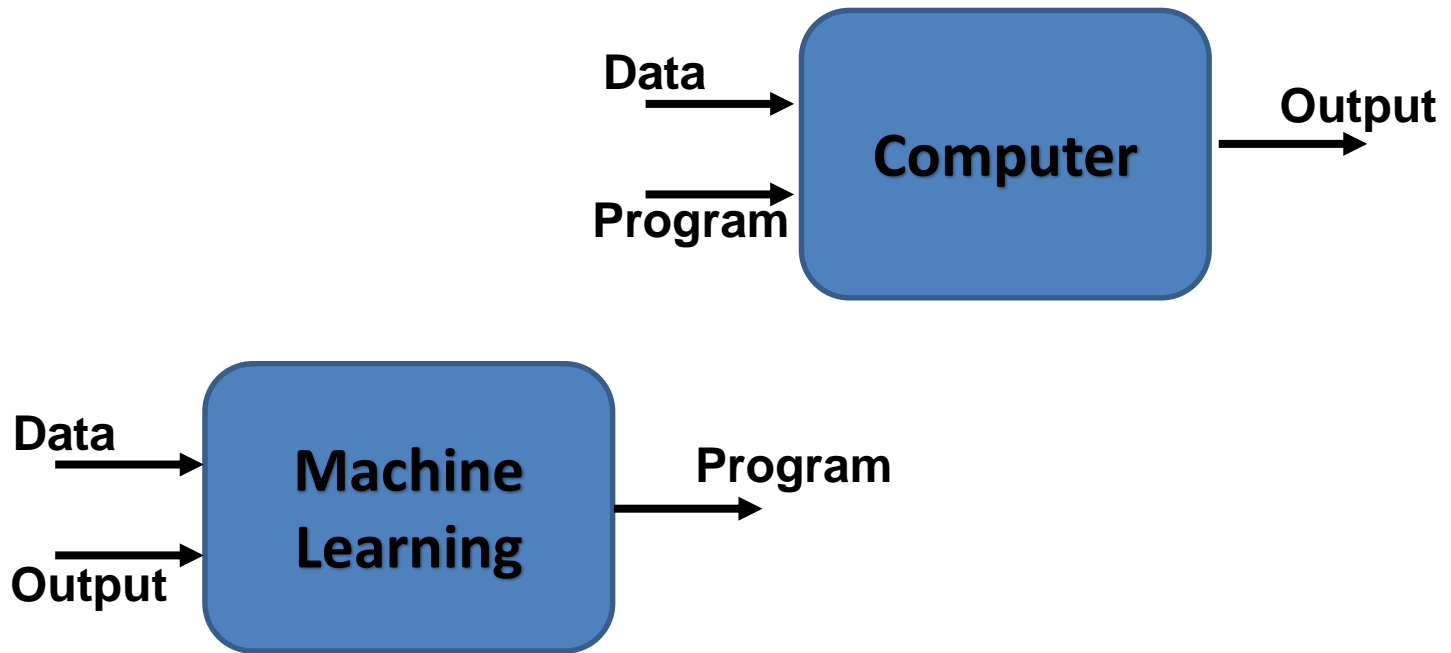
Data Driven Design

- When to use?
 - Difficult to reason about a generic rule that solves the problem
 - Easy to collect examples (with the solution)





What is Machine Learning?

- Automating the Automation





Data Driven Design

- When to use?
 - Difficult to reason about a generic rule that solves the problem
 - Easy to collect examples (with the solution)




Do we need deep learning?

- For simple (linear) classification problems?




Do we need deep learning?

- The Linear Model has its Limits



(a) Linear with outliers




(b) Essentially nonlinear

To address (b) we need something more than linear.

Do we need deep learning?

- Change Your Features Using a Transform

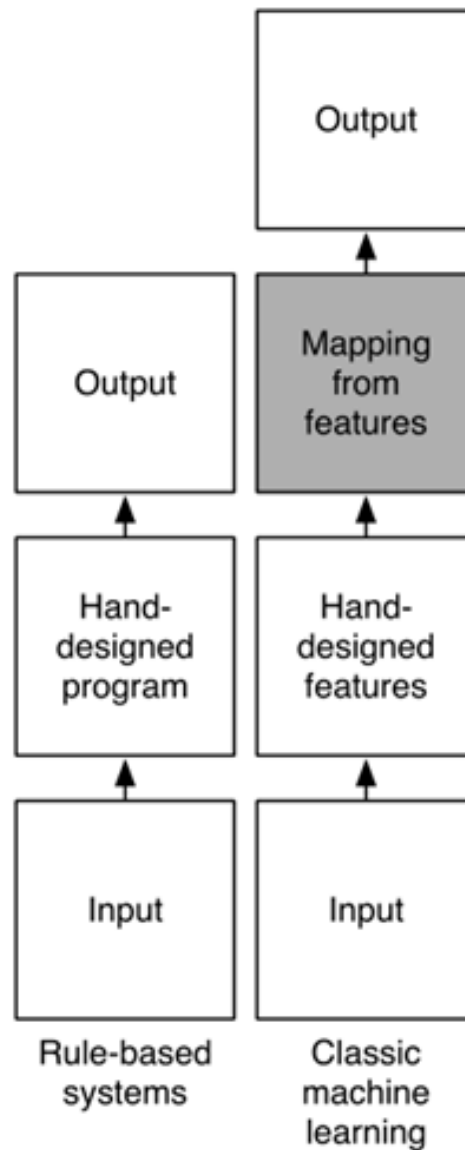


$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}$$



$$\mathbf{z} = \Phi(\mathbf{x}) = \begin{bmatrix} 1 \\ x_1^2 \\ x_2^2 \end{bmatrix} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \Phi_2(\mathbf{x}) \end{bmatrix}$$

Classic Machine Learning




Use the Linear Model?

- First try a linear model – simple, robust and works.
- Algorithms can tolerate error plus you have nonlinear feature transforms.
- Choose a feature transform before seeing the data. Stay simple.
- **Data snooping** is hazardous to your generalization.
- Linear models are fundamental in their own right; they are also the building blocks of many more complex models like support vector machines.
- Nonlinear transforms also apply to regression and logistic regression and ...


Do we need deep learning?

- For complex binary classification problems?







Do we need deep learning?

- For complex binary classification problems?

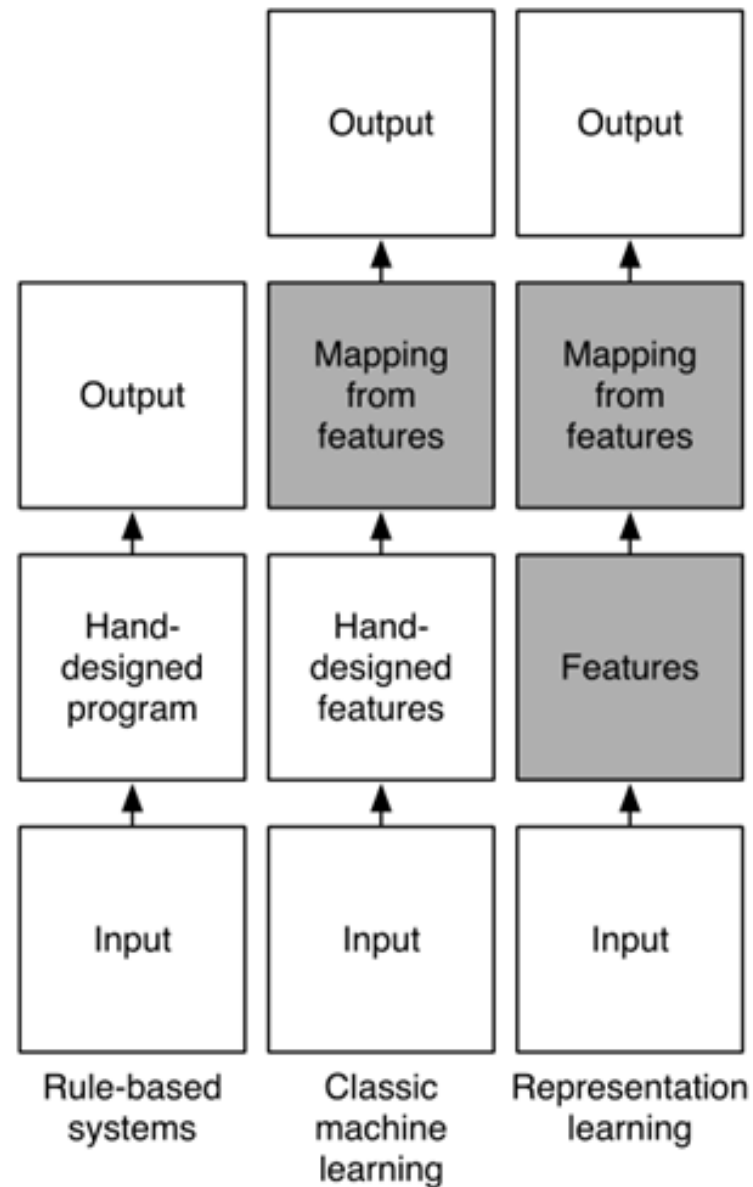


Do we need deep learning?

- For complex binary classification problems?
 - No: **Universal Approximation Property of Feed-forward Networks**



Representation Learning




Do we need deep learning?

- For complex binary classification problems?

Yes

- In any learning task, we have to be concerned with what is feasibly “learnable” in a given representation.
- Using networks with more layers, one can obtain more **compact representations** of the input-output relation.
 - We say that a network is **compact** if it consists of relatively few free parameters (few computational elements) to be learned/tuned during the training phase.
 - For a given number of training points, we expect compact representations to result in better generalization performance.
- For complex tasks, where more complex concepts have to be learned, for example, recognition of a scene in a video recording, language and speech recognition, the underlying functional dependence is of a very complex nature so that we are unable to express it analytically in a simple way.

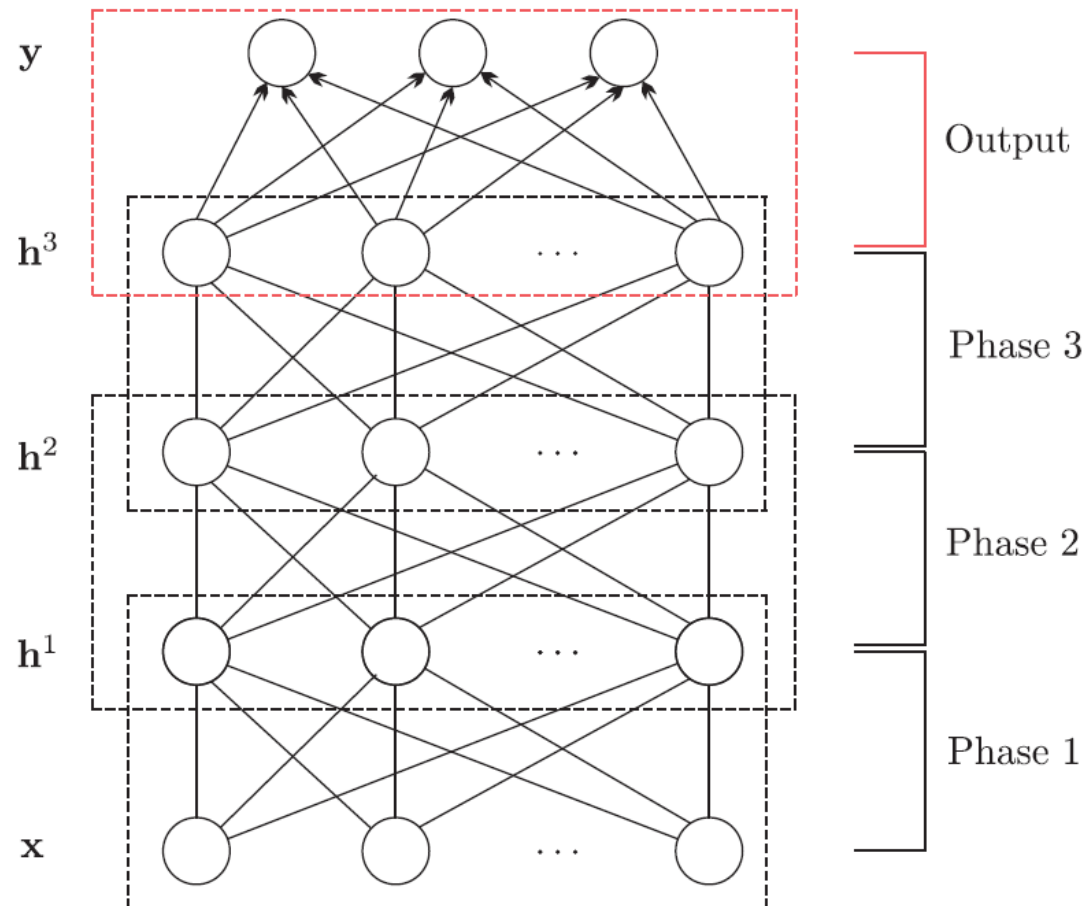


Training Deep Neural Networks

- Top-down Training
 - Training with **backpropagation** can become difficult and often algorithms are stuck in local minima.
 - Is there a training scheme, beyond or complementary to the backpropagation algorithm, to assist the optimization process to settle in a “good” local minimum, by extracting and exploiting more information from the input data?
- Can we **train the representation** without using top-down supervision?

Training Deep Neural Networks





- The main idea is to pre-train each layer, via an **unsupervised** learning algorithm, **one layer at a time**, in a greedy-like rationale.



Training Deep Neural Networks: Different flavours


- Using
Stacked AutoEncoders
(SAE)

Training a SAE

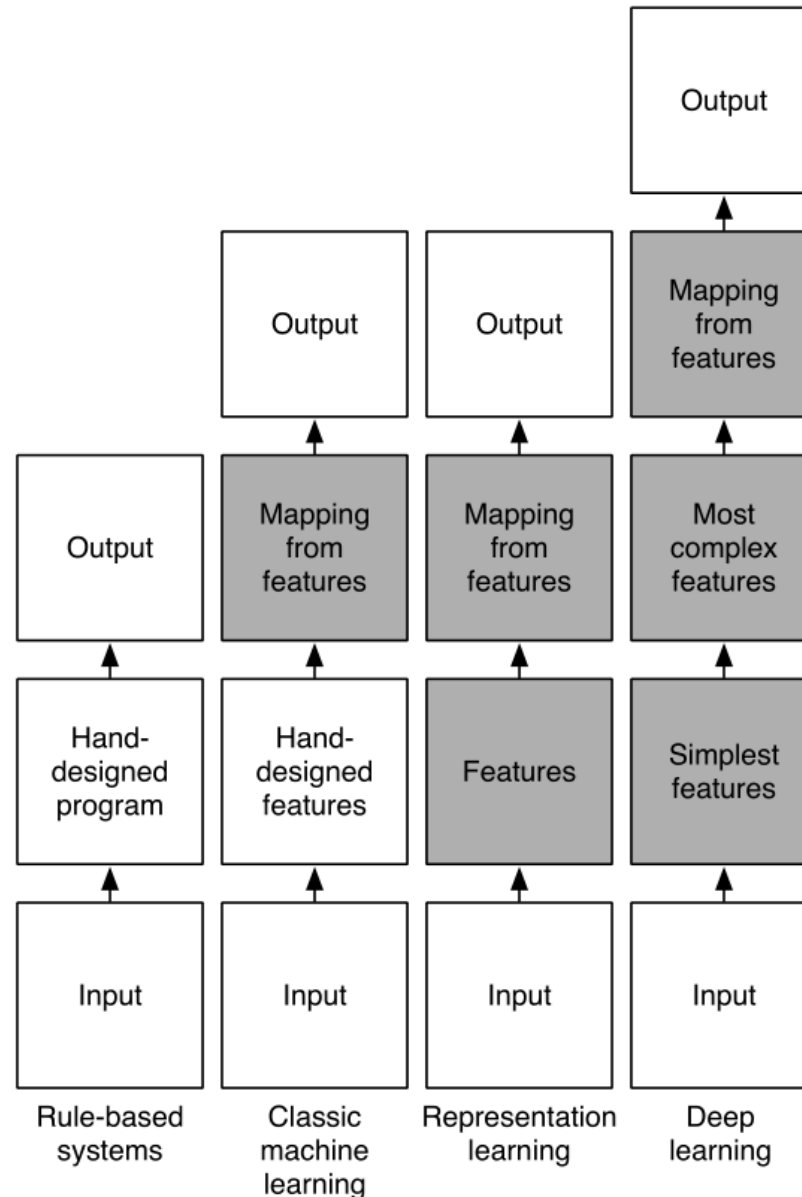


Training Deep Neural Networks: Different flavours

- Training Deep Feed-Forward Networks
 - Step I: Unsupervised Pre-training of Hidden Units
 - Stacked Autoencoders:
 - Phase 1
 - Phase 2
 - Phase 3
 - Step II: Supervised Pre-training of Output Nodes
 - Step III: Fine-Tuning of All Nodes via Supervised Training

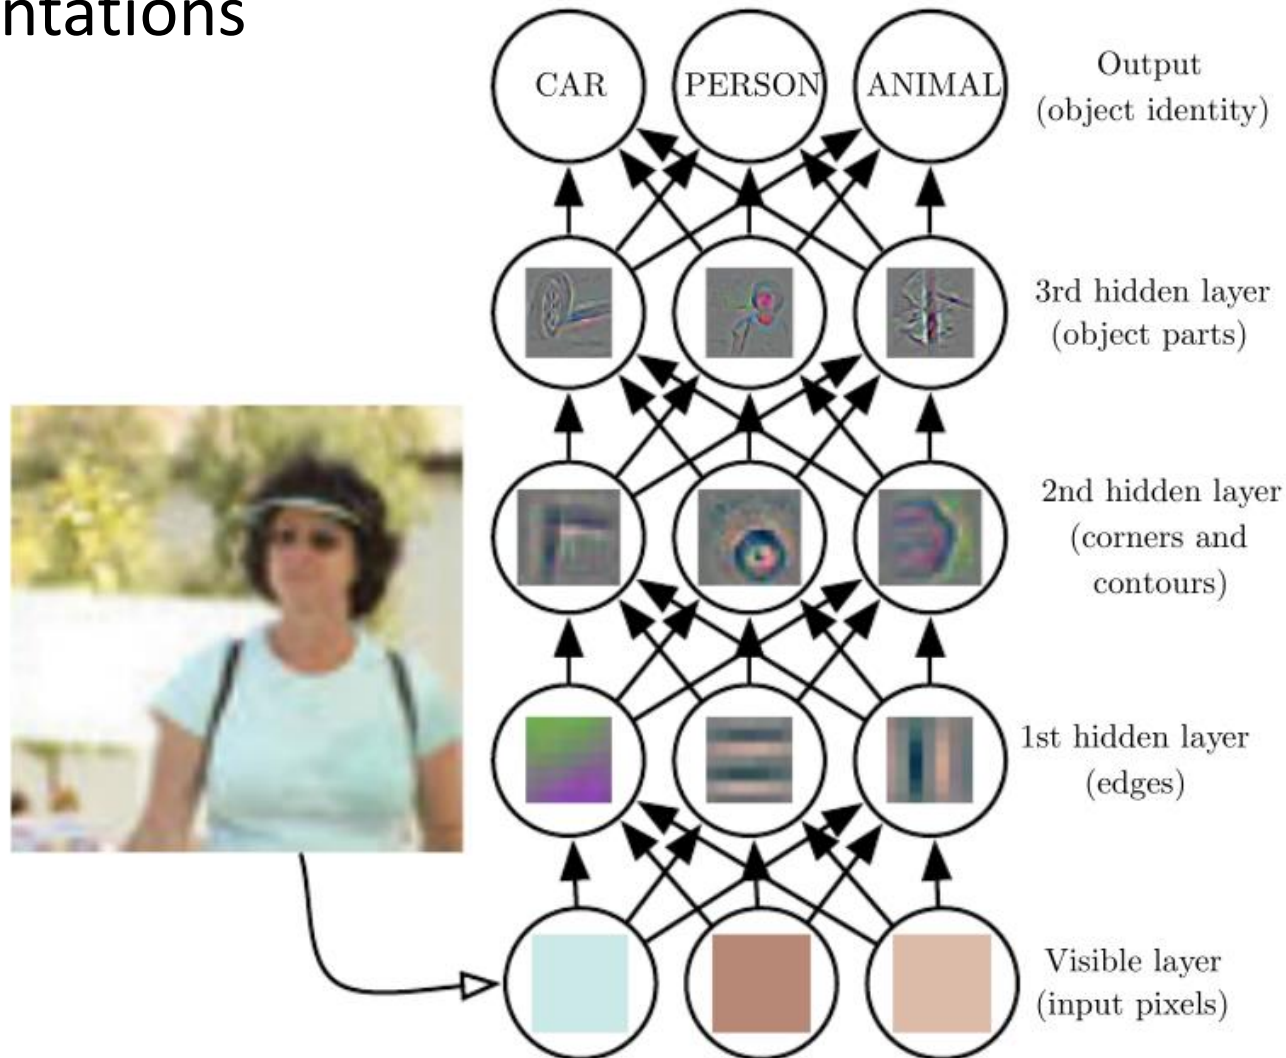


Why is Deep Learning working so well?



Why is Deep Learning working so well?

- Successive model layers learn deeper intermediate representations



Why is Deep Learning working so well?

Four key ingredients

1. Lots & lots of data
2. Very flexible models
3. Enough **computing power**
4. Powerful priors that can defeat the curse of dimensionality

Why Unsupervised (Deep) Representation?

- Potential benefits:
 - Exploit tons of unlabeled data
 - Answer new questions about the variables observed
 - Regularizer – transfer learning – domain adaptation
 - Easier optimization (local training signal)
 - Structured outputs

How do humans generalize from very few examples?

- They transfer knowledge from previous learning:
 - Representations
 - Explanatory factors
- Previous learning from:
unlabeled data + labels for other tasks

Transfer Learning

- Transfer learning (TL) aims to extract knowledge from at least one source task and use it when learning a predictive model for a new target task.
 - The intuition behind this idea is that learning a new task from related tasks should be easier (faster, with better solutions or with less amount of labelled data) than learning the target task in isolation.
 - **transferring data**, namely, strategically including data from the source task in the target dataset
 - adapting knowledge instead of data. This idea is handled by **parameter transfer** approaches, which rely on the idea that individual models for related tasks should share some structure (parameters or hyper-parameters)

Transfer Learning

- In practice, very few people train an entire Convolutional Network from scratch (with random initialization). Instead, it is common to pretrain a ConvNet on a very large dataset
 - ConvNet as fixed feature extractor
 - Fine-tuning the ConvNet

Transfer Learning



- In practice, very few people train an entire Convolutional Network from scratch
 - Example: Face Recognition



- Use VGG Deep Net (trained with millions of faces – but for a **different task**) to extract a feature vector by removing the output layer
- Use euclidean distance (or variations) for matching
 - Improves over state of the art approaches in several settings.


Transfer Learning

- Cross Sensor Adaptation for periocular recognition




Multi-Task Learning

- Generalizing better to new tasks (tens of thousands!) is crucial to approach AI
- Example: speech recognition, sharing across multiple languages
- Deep architectures learn good intermediate representations that can be shared across tasks




E.g. dictionary, with intermediate concepts re-used across many definitions



Multimodal Deep Learning






Design of a Classifier



Design of a Classifier




Design of a Classifier

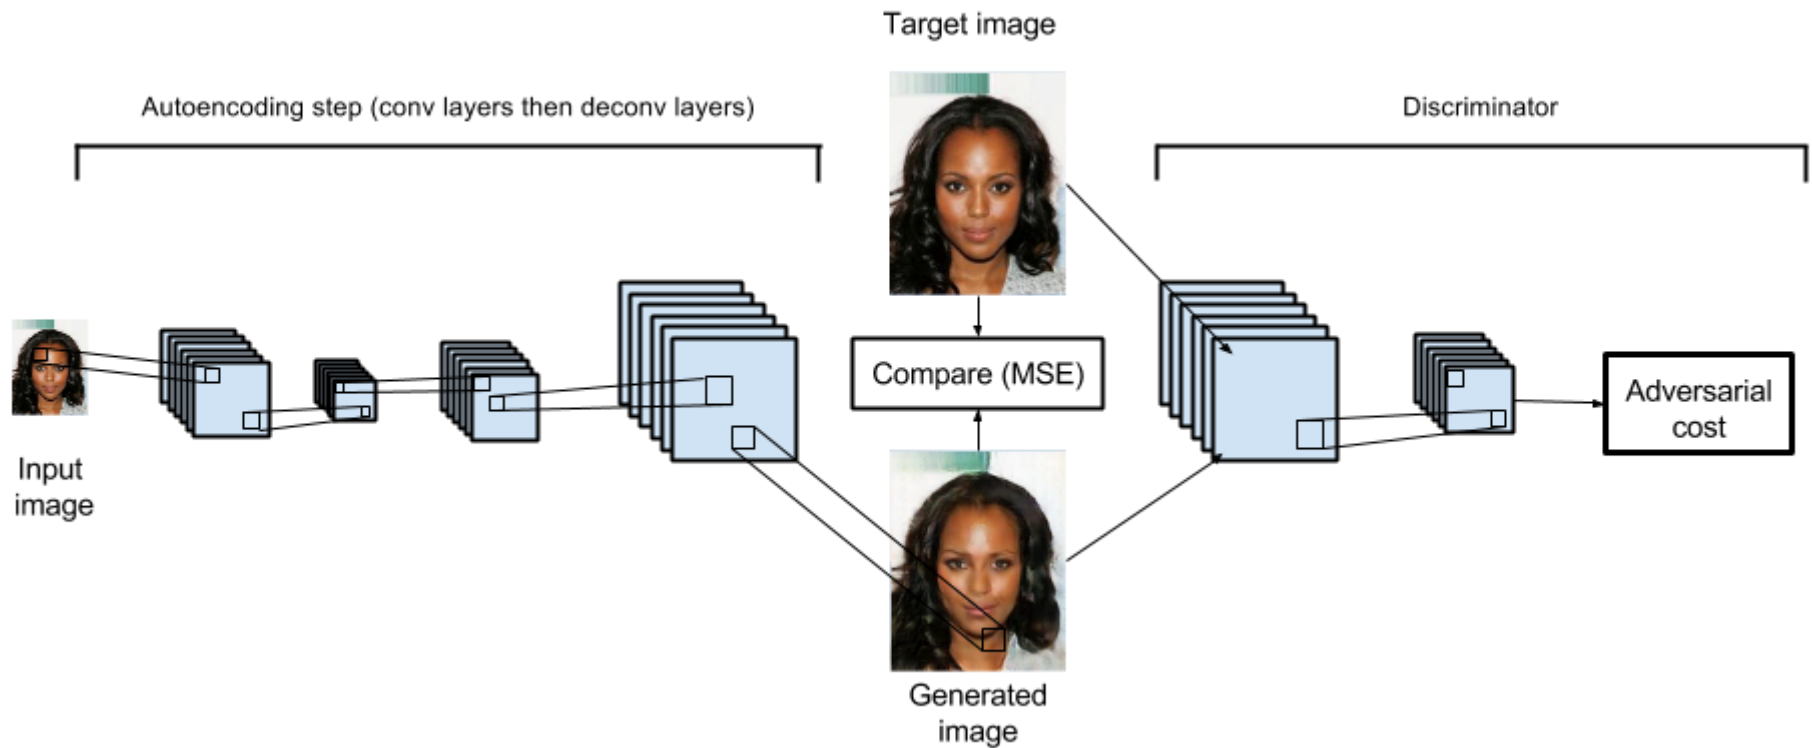


Deep Belief Networks

- We focused our discussion on deep learning on multilayer perceptrons for supervised learning.
 - Focus on the information flow in the feed-forward or bottom-up direction.
- **Can we start from the last layer, corresponding to the most abstract representation, and follow a **top-down path** with the new goal of generating data?**
 - Studies suggest that such **top-down connections** exist in our visual system to generate lower level features of images starting from higher level representations.
 - Important for disambiguating effect on the interpretation of local image regions by providing contextual prior information from previous frames



Adversarial Neural Network



Adversarial Neural Network

Noise $\sim N(0,1)$



Generative
Model



Ideas

- Ideally: the Master Algorithm
 - Given
 - data
 - learning bias / prior
 - Does the job
- Deep Models: a step in the right direction?
 - A lot of parameters
 - Difficult to guide the learning in the ‘right direction’
 - Difficult to avoid overfitting
 - Presently we are just overfitting to larger datasets
 - Vulnerable to **adversarial examples**

References (slides/images/technical info)

- <http://cs231n.stanford.edu/>
 - <http://cs231n.github.io/convolutional-networks/>
 - <http://cs231n.github.io/transfer-learning/>
- Machine Learning A Bayesian and Optimization Perspective
- Bengio's webpage
 - http://www.iro.umontreal.ca/~bengioy/yoshua_en/index.html
- Deep Learning, Ian Goodfellow, Yoshua Bengio, Aaron Courville
- Jaime Cardoso's Webpage
 - <http://www.inescporto.pt/~jsc/>

THANK YOU